

CHASR: Causal Hallucination Attribution for Extreme Super-Resolution

Quang-Vinh Dang

British University Vietnam, Hung Yen, Vietnam

Abstract

Extreme super-resolution (xSR, $8\times/16\times$ and beyond) is severely ill-posed: the diffusion priors that make plausible upscaling possible also make some hallucinated, plausible-but-false detail unavoidable. Existing work measures how much a super-resolved image hallucinates but not why, which limits its usefulness for forensic and task-driven applications needing to know whether a detail can be trusted. We introduce CHASR, a causal attribution module decomposing hallucination risk at the pixel level into four distinguishable causes: inherent ill-posedness, over-reliance on the generative prior, degradation-model mismatch, and distribution shift from the backbone’s training data. CHASR wraps a frozen chained diffusion backbone with four lightweight signal estimators and a small trained fusion head producing a per-pixel five-way cause classification plus a calibrated reliability score. We construct xSR-CausalBench, a synthetic benchmark injecting each cause under controlled conditions with weak per-pixel ground truth, extending an existing two-way controllable-hallucination construction to four causes. A first end-to-end validation at $16\times$ on real photographs (40 training, 20 held-out images) shows the fusion head reaches 48.4% five-way pixel attribution accuracy, above a uniform-random baseline (20%) but only 0.7 points above the benchmark’s empirical majority-class baseline (47.7%), a gap attributed to unweighted training on an imbalanced benchmark at pilot scale rather than the architecture. A disentanglement analysis confirms an an-

Email address: `vinh.dq4@buv.edu.vn` (Quang-Vinh Dang)

anticipated factor-entanglement risk: the null-space-uncertainty and prior-reliance signals correlate at $r = 0.267$, consistent with severe ill-posedness driving prior reliance. We report this pilot honestly, including where the numbers are weak. The central contribution is the causal taxonomy and attribution framework, not fidelity leaderboard performance.

Keywords: extreme super-resolution, hallucination attribution, diffusion models, causal analysis, null-space uncertainty, image forensics

1. Introduction

A security camera captures a license plate at a distance where the plate occupies perhaps a dozen pixels. A satellite image resolves a rooftop as a handful of gray values. A decades-old photograph, scanned at low resolution, is the only surviving record of a document. In each case, someone eventually asks a diffusion-based super-resolution system to recover detail that the sensor never captured, and in each case the system will comply: modern extreme super-resolution (xSR) methods, operating at $8\times$ or $16\times$ upscaling factors, now produce images that look convincingly sharp. The trouble is that “looks convincingly sharp” and “is actually correct” have become only loosely related properties. At these factors the information-theoretic content of the input is a small fraction of what the output claims to show, and the generative prior inside the model is doing most of the work of deciding what the missing pixels should be. When it decides wrong, the result is hallucination: detail that is plausible, consistent with the model’s training distribution, and false.

The severity of this problem scales directly with how the output will be used. A hallucinated license-plate character that happens to look legible is not a cosmetic defect if that character becomes evidence. A hallucinated texture on a satellite rooftop is not a rendering artifact if an analyst reports it as a structural feature. This is precisely the concern that motivates growing interest in extreme super-resolution’s forensic and task-driven applications, and it explains why the special issue this paper targets asks explicitly for evaluation metrics that go

beyond PSNR and SSIM: a fidelity score computed against a ground-truth image the deployment scenario does not have access to cannot tell an analyst whether
25 the ten pixels they are looking at right now can be trusted.

A handful of recent papers move past fidelity to measure hallucination directly. HalluGen [1] introduces a controllable benchmark that separates intrinsic hallucination (violations of data consistency) from extrinsic hallucination (wrong content filled into the model’s null space) and scores restorations against
30 this two-way taxonomy. The Hallucination Score [2] takes a complementary approach, using a multimodal large language model to produce a single scalar severity score that can double as a differentiable fine-tuning signal. Both are genuine advances, and both share a structural limitation for the use cases this paper is concerned with: they tell a user *how much* a region hallucinates, or
35 at best which of two broad categories the hallucination falls into, but not *why*. An analyst deciding whether to trust a specific license-plate character needs more than a severity number. They need to know whether the character is uncertain because the optics genuinely destroyed that information (in which case no model, however good, could have recovered it correctly), because the model
40 ignored usable evidence in favor of its own prior (a model defect, potentially fixable), because the model was applied to a degradation it was never trained to handle (an operating-conditions mismatch), or because the content itself is unlike anything the model has seen (an out-of-distribution failure). These are four distinguishable failure modes with four different practical responses, and
45 no existing method attributes hallucination to a specific one of them.

This paper introduces CHASR (Causal Hallucination Attribution for Super-Resolution), a module that sits on top of a frozen extreme super-resolution backbone and produces, for every pixel of the output, both a reliability score and a classification into one of four causes or a fifth “reliable” class. The design
50 commits to four causal signals that map directly onto the four failure modes above: null-space uncertainty, measured as the pixelwise variance across multiple degradation-consistent stochastic samples from the backbone; prior-reliance, measured as the relative sensitivity of the output to the low-resolution evidence

versus to the model’s random seed; degradation-model mismatch, measured by
55 re-estimating the input’s blur and noise parameters with an independent blind
kernel estimator and comparing them against the backbone’s assumed training
distribution; and distribution shift, measured as the nearest-neighbor distance
from the output’s frozen self-supervised features to a bank of the backbone’s
training-distribution features. A small fusion head, the only trained component
60 in the entire pipeline, learns to combine these four raw signals with a shallow
image feature into a calibrated per-pixel attribution.

Training and evaluating this fusion head requires per-pixel ground truth for
which cause produced a given hallucination, and no such labels exist for real
images because the true cause of a real hallucination is, definitionally, unobserv-
65 able. We address this by constructing xSR-CausalBench, a synthetic benchmark
that extends HalluGen’s controllable-hallucination construction from two causes
to four: four distinct procedures inject ill-posedness, prior-reliance, degradation-
mismatch, or distribution-shift dominance into an otherwise normal image un-
der experimentally controlled conditions, producing a weak per-pixel label mask
70 alongside each synthesized low-resolution and high-resolution pair.

The causes this taxonomy separates are not independent in practice, and
pretending otherwise would misrepresent what the framework can promise. A
region with extreme ill-posedness naturally invites more prior reliance, because
there is less real evidence available for the model to weigh against its prior in
75 the first place. Section 5 reports this correlation directly, measured from the
pipeline’s own output rather than assumed, because the alternative, quietly
hoping the four signals turn out to be independent and finding out otherwise at
review time, serves nobody.

This paper’s contribution is fourfold: a four-way causal taxonomy for ex-
80 treme super-resolution hallucination, extending the two-way intrinsic/extrinsic
split in the closest prior work; the CHASR architecture that instantiates the
taxonomy as four computable signals plus a learned fusion head, built on a pub-
licly reproducible diffusion backbone; xSR-CausalBench, a synthetic benchmark
that generates per-pixel weak ground truth for all four causes; and an honestly

85 reported pilot validation, at real $16\times$ upscaling on natural photographs, that demonstrates the pipeline is trainable and produces above-chance, structurally sensible attribution, while being explicit about what remains to scale before the numbers reported here should be read as a final answer rather than a first one.

2. Related Work

90 2.1. *Blind and diffusion-based super-resolution*

Classical blind super-resolution treats the degradation kernel as an unknown to be estimated jointly with, or prior to, the restoration itself. KernelGAN [3] estimates a per-image blur kernel using an internal GAN trained purely on patch recurrence within the low-resolution image, requiring no external training data. 95 MANet [4] instead learns a mutual affine network that estimates spatially variant kernels directly, relaxing the common assumption of a single global kernel per image. DASR [5] sidesteps explicit kernel parameterization entirely, learning an unsupervised degradation representation via contrastive learning and conditioning the restoration network on it directly. This lineage of blind kernel 100 estimation is a component CHASR reuses directly: Signal 3 (degradation-model mismatch) is computed by an independent blind estimator trained on both the backbone’s assumed degradation pool and an out-of-distribution pool, then compared against the backbone’s training assumptions, rather than proposed here as a new estimation method.

105 Diffusion models have displaced GAN-based approaches as the dominant paradigm for photorealistic super-resolution at moderate factors. StableSR [6] adapts a pretrained text-to-image diffusion prior for real-world restoration through a time-aware encoder and a controllable feature-wrapping module. SeeSR [7] conditions the diffusion process on semantic tags extracted from the degraded 110 input, improving fidelity to the true image content rather than merely to a plausible one. PASD [8] introduces pixel-aware cross-attention so that the diffusion process attends explicitly to which regions the input actually constrains. Chain-of-Zoom [9] pushes this diffusion-prior paradigm to substantially more extreme

factors by reformulating extreme super-resolution as an autoregressive sequence
of moderate zoom-in steps guided by vision-language model text prompts, each
115 step handled by an off-the-shelf $4\times$ diffusion model rather than a single model
retrained for the full factor. CHASR’s backbone follows this same chained-hop
principle for reaching $16\times$ from a $4\times$ checkpoint, though for reproducibility on
modest hardware it wraps the publicly downloadable `stable-diffusion-x4-u`
120 `pscaler` checkpoint directly rather than a VLM-guided variant, trading some of
Chain-of-Zoom’s content-awareness for a fully open, easily reproduced pipeline.
Real-ESRGAN [10] remains the standard GAN-based reference point for real-
world super-resolution trained entirely on synthetic degradations, and Section 4
reports it as a real fidelity baseline, chained to the same $16\times$ factor as CHASR’s
125 own backbone, rather than only citing it here. Two concurrent lines of work cou-
ple diffusion-based super-resolution directly to an uncertainty signal rather than
treating uncertainty as an afterthought: UGDiff [11] estimates per-latent-feature
reconstruction uncertainty and uses it to steer high-frequency detail restoration
toward high-uncertainty regions while leaving confident regions untouched, and
130 QUSR [12] adapts injected diffusion noise strength to local uncertainty along-
side a multimodal-LLM-based quality prior, under unknown, spatially varying
degradation. Both use uncertainty as a single scalar steering signal for where to
generate more aggressively; CHASR’s Signal 1 measures a structurally similar
null-space quantity but treats it as one of four distinguishable causes to attribute
135 after generation rather than a knob to steer generation with, a diagnostic rather
than a generative use of the same underlying idea.

2.2. Measuring hallucination in generative restoration

HalluGen [1] is the closest prior work to CHASR’s evaluation methodology.
It separates hallucination into intrinsic (violating data consistency with the
140 observed degradation) and extrinsic (wrong content filled into the null space
where data consistency alone does not constrain the answer) categories, con-
structs a synthetic controllable-hallucination benchmark via diffusion posterior
sampling, and proposes a patch-level severity metric. Its own benchmark is built

from low-field MRI restoration rather than natural photographs, so CHASR’s
 145 xSR-CausalBench (Section 3.5) is a cross-domain methodological descendant
 rather than a direct reuse: it ports the controllable-injection idea from medi-
 cal to natural-image extreme super-resolution, extending it from two categories
 to four and from a severity score to a per-pixel causal label. The Halluci-
 nation Score [2] takes a different but complementary route: rather than de-
 150 composing hallucination by cause, it uses a multimodal large language model
 to judge two properties, implausibility under any degradation and semantic
 deviation from what the input evidence supports, and combines them into a
 scalar that doubles as a reward signal for fine-tuning. Both works measure
 severity or coarse category; neither attributes a specific hallucinated region
 155 to one of several distinguishable causes, which is the gap this paper targets
 directly. The same detection-versus-attribution gap has emerged concurrently
 outside super-resolution: Vo et al. [13] decompose vision-language model failures
 on knowledge-intensive visual question answering into recognition-stage versus
 knowledge-stage causes rather than a single failure score, enabling different fixes
 160 (image repair, entity support, question rewriting) for different causes, the same
 why-before-what argument this paper makes for extreme super-resolution. In a
 related line of the present author’s own work, addressing a structurally similar
 problem in a different modality, Dang et al. [14] intercept and correct halluci-
 nated tokens during vision-language model generation via a dynamically con-
 165 structed, externally-verified knowledge graph rather than a post-hoc severity
 score, underscoring that mechanism-aware, causal intervention is an emerging
 theme across generative-AI hallucination research generally, not a concern spe-
 cific to image restoration.

At a more general level, work on diffusion models outside restoration has
 170 begun explaining hallucination as an out-of-distribution phenomenon: Aithal et
 al. [15] show that diffusion models can hallucinate through mode interpolation,
 generating samples in regions between distinct modes of the training distribu-
 tion that the model never actually observed, which is a mechanistic account
 of why unfamiliar content produces confident-looking but wrong output. Sig-

nal 4 in this paper’s method operationalizes the restoration-specific analogue of that same intuition, measuring an output’s feature-space distance from the backbone’s training distribution directly, rather than proposing a new theory of why distribution shift causes hallucination.

Range-null space decomposition, formalized for diffusion-based zero-shot inverse problems by DDNM [16], provides the mathematical grounding for CHASR’s Signal 1. DDNM shows that for a known linear degradation operator, a diffusion sampling trajectory can be projected at every step so that the range-space component exactly matches the observation while the diffusion prior is left free to fill the null space, and demonstrates this on a range of restoration tasks under the assumption that the degradation operator is known in advance. CHASR uses the same range-null projection, applied to the average-pooling operator that governs its synthetic degradation pipeline, to guarantee that every stochastic sample the backbone produces is exactly degradation-consistent, which is what makes the sample-to-sample variance in Signal 1 attributable to null-space uncertainty specifically rather than to a data-consistency violation. Where DDNM assumes a known degradation and does not address attribution, CHASR assumes a partially unknown degradation (estimated blindly for Signal 3) and uses the resulting variance as one input to a diagnostic rather than as the whole restoration method.

2.3. Task-driven super-resolution and this special issue

The special issue this paper targets is organized in part around task-driven reliability, and closely related work on embedding-similarity-guided license-plate super-resolution [17] exemplifies the concern directly: a super-resolved license plate is only useful if the specific characters it renders are correct, not merely if the overall image looks sharp. The concurrent XLPSR Grand Challenge at ICIP 2026, whose finalist entries include ensemble approaches combining super-resolution backbones with scene-text recognition voting [18], reflects the same practical stakes at competition scale: entries are scored on downstream recognition accuracy, not perceptual quality. CHASR’s downstream case study,

205 correlating flagged high-mismatch or high-distribution-shift regions with actual optical character recognition failure on text-patch crops rather than treating task-driven evaluation as secondary to fidelity, is implemented as a reusable module (Section 5) but not yet exercised against real recognition failures in this pilot.

210 3. Method

CHASR is a diagnostic module, not a new super-resolution architecture. It wraps a frozen backbone and adds four causal signal estimators plus a small trained fusion head, so that the entire training cost of the system is confined to the fusion head while the backbone’s own super-resolution quality is left exactly as published. Figure 1 shows the full pipeline: the frozen backbone (Section 3.2) produces the super-resolved output and the raw material each signal estimator (Section 3.3) needs, and the only trained component, the fusion head (Section 3.4), combines the four resulting signal maps into a per-pixel cause classification and reliability score.

220 3.1. Problem formulation

Let $x \in [0, 1]^{H \times W \times 3}$ denote a high-resolution image and A_r a linear degradation operator that block-average-pools by factor r , so for every non-overlapping $r \times r$ block B ,

$$(A_r x)_B = \frac{1}{r^2} \sum_{i \in B} x_i. \quad (1)$$

The synthetic low-resolution observation is $y = A_r(k * x)$, where k is a blur kernel applied before pooling. Every blur kernel used in this paper, whether it produces an in-distribution or out-of-distribution degradation, is drawn from the same general anisotropic Gaussian family,

$$k(u, v; \sigma_x, \sigma_y, \theta) \propto \exp\left(-\frac{u_\theta^2}{2\sigma_x^2} - \frac{v_\theta^2}{2\sigma_y^2}\right), \quad \begin{bmatrix} u_\theta \\ v_\theta \end{bmatrix} = \begin{bmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}, \quad (2)$$




Figure 1: CHASR pipeline. Blue: frozen backbone. Green: the four independent signal estimators, each reading the backbone’s output and/or intermediate state (Section 3.3). Orange: the fusion head, the only trained component, which stacks the four signal maps with a shallow luminance feature and predicts a per-pixel cause and reliability.

normalized to sum to 1 over its support; the standard (in-distribution) pool fixes $\sigma_x = \sigma_y$ and $\theta = 0$ (isotropic), while the mismatch pool draws σ_x, σ_y independently from a disjoint, larger range and θ from a nonzero range (Section 3.5), making Equation 2 the shared parametric family Signal 3’s estimator regresses $(\hat{\sigma}_x, \hat{\sigma}_y, \hat{\theta})$ against.

Extreme super-resolution seeks $\hat{x} \approx x$ from y alone at $r = 16$, a severely underdetermined inverse problem: A_{16} discards 255/256 of the pixel-level information in x , so any $\hat{x}$ satisfying $A_{16}(\hat{x}) = y$ is a valid data-consistent completion regardless of whether it matches the true x . CHASR does not attempt to resolve this ill-posedness; it attributes it. Given the frozen backbone’s output $\hat{x}$ and a fixed taxonomy of causes $\mathcal{C} = \{\text{ILL_POSED}, \text{PRIOR_RELIANCE}, \text{MISMATCH}, \text{DIST_SHIFT}, \text{RELIABLE}\}$,

the task this paper addresses is learning a per-pixel classifier

$$\hat{c}(p) = \arg \max_{c \in \mathcal{C}} f_{\theta}(S_1(p), S_2(p), S_3(p), S_4(p), L(p))_c, \quad (3)$$

240 where $S_1, \dots, S_4$ are the four causal signals defined in Section 3.3, L is a shallow luminance feature, and f_{θ} is the fusion head of Section 3.4 — the only component in Equation 3 with learned parameters θ ; every signal on its right-hand side is computed from a frozen or independently pretrained model.

3.2. Frozen backbone

245 The backbone reaches $16\times$ from 64×64 low-resolution patches through two chained $4\times$ diffusion hops, using the publicly downloadable checkpoint **stable-diffusion-x4-upscaler** (hosted on the Hugging Face Hub under the **stabilityai** organization) at each hop for full reproducibility on a single consumer GPU. The first hop upscales the full 64×64 patch to 256×256 . The second
250 hop is applied per-tile, splitting the 256×256 intermediate into four 128×128 tiles, each independently upscaled to 512×512 and stitched back into the final 1024×1024 output, which keeps peak memory within an 8 GB budget at the cost of a faint seam artifact at tile boundaries that does not affect the causal signals computed here. After every hop, the output is projected via the DDNM
255 range-null decomposition [16] against the true low-resolution observation at that hop’s scale. For a linear degradation operator A_r , DDNM’s general projection replaces a raw diffusion estimate $\hat{x}_{\theta}$ with

$$\hat{x}' = A_r^{\dagger} y + (I - A_r^{\dagger} A_r) \hat{x}_{\theta}, \quad (4)$$

where $A_r^{\dagger}$ is the Moore–Penrose pseudo-inverse of A_r : the first term fixes the range space to exactly match the observation, and the second leaves the null
260 space, everything A_r cannot see, free for the diffusion prior to fill. For the block-average-pooling operator of Equation 1, A_r has orthogonal rows ($A_r A_r^{\top} = r^{-2} I$), so $A_r^{\dagger} y$ broadcasts each observed value y_B uniformly across its source block, and $A_r^{\dagger} A_r \hat{x}_{\theta}$ replaces every pixel with its own block mean. Equation 4 then reduces,

block by block, to

$$\hat{x}'_i = \hat{x}_{\theta,i} - \frac{1}{r^2} \sum_{j \in B} \hat{x}_{\theta,j} + y_B, \quad \forall i \in B, \quad (5)$$

265 the closed-form, per-block correction actually implemented: each pixel is shifted
by the difference between its block’s true observed mean and the diffusion out-
put’s own block mean. This guarantees every sample the backbone returns ex-
actly satisfies $A_r(\hat{x}') = y$, at every hop, which is what makes the variance across
independently sampled outputs attributable specifically to null-space content
270 rather than to the model occasionally violating the observation. Algorithm 1
states the full chained procedure, applying Equation 5 at each hop’s own scale.

Algorithm 1: Chained, DDNM-projected 16× sampling

Input : Low-resolution patch $y \in \mathbb{R}^{64 \times 64 \times 3}$, sample count K

Output: K degradation-consistent 16× super-resolved samples

$\{\hat{x}_1, \dots, \hat{x}_K\}$

```

1 for  $k \leftarrow 1$  to  $K$  do
2    $z_1 \leftarrow \text{DiffusionHop}_1(y, \text{seed} = k)$  ;           //  $64 \times 64 \rightarrow 256 \times 256$ 
3    $z_1 \leftarrow \text{DDNM\_project}(z_1, y, \text{scale} = 4)$  ; // exact consistency at
   hop 1's scale
4   for  $\text{tile} \in \text{split}(z_1, 4)$  do
5      $t \leftarrow \text{DiffusionHop}_2(\text{tile}, \text{seed} = k)$  ;    //  $128 \times 128 \rightarrow 512 \times 512$ 
6      $t \leftarrow \text{DDNM\_project}(t, \text{tile}, \text{scale} = 4)$ ;
7      $\text{stitch}(t)$  into  $z_2 \in \mathbb{R}^{1024 \times 1024 \times 3}$ ;
8   end
9    $\hat{x}_k \leftarrow \text{DDNM\_project}(z_2, y, \text{scale} = 16)$  ; // final consistency
   check against  $y$  directly
10 end
11 return  $\{\hat{x}_1, \dots, \hat{x}_K\}$ 

```

3.3. Four causal signals

Signal 1: null-space uncertainty. The backbone is sampled K times with different random seeds, each sample $\hat{x}_1, \dots, \hat{x}_K$ independently satisfying Equation 5. The pixelwise, channel-averaged unbiased sample variance across the K samples,

$$\bar{S}_1(p) = \frac{1}{3} \sum_{ch=1}^3 \frac{1}{K-1} \sum_{k=1}^K (\hat{x}_k(p, ch) - \bar{x}(p, ch))^2, \quad \bar{x}(p, ch) = \frac{1}{K} \sum_{k=1}^K \hat{x}_k(p, ch), \quad (6)$$

is then rescaled to $[0, 1]$ against the theoretical maximum unbiased two-sample variance a $[0, 1]$ -bounded pixel value can attain (achieved when $\hat{x}_1(p) = 0, \hat{x}_2(p) = 1$, giving $\bar{S}_1^{\max} = 0.5$),

$$S_1(p) = \min(1, \bar{S}_1(p) / \bar{S}_1^{\max}). \quad (7)$$

Because data consistency is fixed by construction, the samples can only differ in their null-space content, so high variance at a pixel means the ill-posedness at that pixel is large enough that the model genuinely could not have determined a unique answer from the evidence, independent of whether the model itself is well-calibrated.

Signal 2: generative-prior over-reliance. This signal asks a different question: even where evidence is available, did the model use it? For the frozen diffusion backbone, whose stochasticity is a discrete random seed rather than a continuous latent, this is estimated by finite differences at the first (256×256) hop: the raw, pre-projection output’s sensitivity to a small perturbation of the low-resolution evidence (g_{evidence}) is compared against its sensitivity to a full reseed at fixed evidence (g_{prior}),

$$S_2(p) = \frac{g_{\text{prior}}(p)}{g_{\text{prior}}(p) + g_{\text{evidence}}(p) + \varepsilon}. \quad (8)$$

Both terms are compared as raw output-magnitude changes on the same units, rather than as normalized derivatives, because normalizing only the continuous perturbation by its step size while leaving the discrete reseed unnormalized would bias the ratio toward zero regardless of the model’s actual behavior.

A general finite-difference formulation for a continuously perturbable forward function and latent is also implemented as a standalone, independently tested utility for future backbones with a true continuous latent, but is not used in the real pipeline for exactly this reason.

300 **Signal 3: degradation-model mismatch.** An independent, lightweight convolutional network is trained to regress a blur kernel’s spatial parameters (anisotropic sigma $(\hat{\sigma}_x, \hat{\sigma}_y)$, rotation angle $\hat{\theta}$) and additive noise level directly from a 64×64 patch, trained on a 50/50 mixture of the backbone’s assumed in-distribution degradation pool (upper sigma bound $\sigma_{\max}$, zero rotation) and
 305 an out-of-distribution pool with a disjoint, larger sigma range and non-zero rotation. At inference, this estimator re-examines the actual input patch and scores how far the estimate falls outside the assumed pool,

$$S_3 = \min\left(1, \frac{\max(0, \max(\hat{\sigma}_x, \hat{\sigma}_y) - \sigma_{\max}) + |\hat{\theta}|}{d_{\max}}\right), \quad (9)$$

normalized by the out-of-distribution pool’s own worst-case raw distance $d_{\max}$ so the signal saturates at 1.0 only at genuinely extreme mismatch rather than
 310 prematurely. S_3 is a single scalar per input patch, broadcast uniformly to every pixel, since the degradation kernel is assumed spatially constant within one patch. A high value here means the backbone is being asked to restore an image degraded in a way it was never exposed to during training, independent of whether the content itself is familiar.

315 **Signal 4: distribution shift.** A frozen, self-supervised DINOv2 encoder [19] extracts a feature vector f from the super-resolved output, and its k -nearest-neighbor distance $d_{k\text{NN}}(f, \mathcal{B})$ to a reference bank $\mathcal{B}$ of features extracted once from the backbone’s training-distribution images is computed. Because a sparse reference bank in a high-dimensional feature space suffers a curse-of-dimensionality
 320 effect at $k > 1$, where a near-duplicate query’s second through k -th neighbors are effectively generic points at a roughly constant distance that washes out fine-grained separation, the raw distance is rescaled by the bank’s own internal

typical k -nearest-neighbor distance $s(\mathcal{B})$ among its own members,

$$S_4 = \frac{d_{\text{kNN}}(f, \mathcal{B})}{d_{\text{kNN}}(f, \mathcal{B}) + s(\mathcal{B})}, \quad (10)$$

producing a self-calibrated score in $[0, 1]$ without assuming a fixed feature-space
 325 magnitude in advance. Like S_3 , S_4 is computed once per output image and
 broadcast to every pixel.

Algorithm 2 collects the four signal computations (Equations 6–10) in one
 place, making explicit which quantities each signal shares (all four read from
 the same K -sample set of Algorithm 1, and S_2 additionally requires two extra
 330 single-hop calls beyond it) and which are independently pretrained rather than
 computed fresh per pixel (S_3 ’s kernel estimator E_θ and S_4 ’s DINOv2 encoder
 ϕ are both frozen, queried here, not trained here).

Algorithm 2: Four-signal computation

Input : LR patch y , samples $\{\hat{x}_1, \dots, \hat{x}_K\}$ from Algorithm 1, kernel
 estimator E_θ , DINOv2 encoder ϕ , reference bank $\mathcal{B}$

Output: Signal stack $[S_1, S_2, S_3, S_4] \in [0, 1]^{4 \times 1024 \times 1024}$

- 1 $S_1 \leftarrow$ Eq. 6, then normalized via Eq. 7;
- 2 $h_0 \leftarrow \text{DiffusionHop}_1(y, \text{seed} = 0)$;
- 3 $h_{\text{ev}} \leftarrow \text{DiffusionHop}_1(y + \epsilon \cdot \mathcal{N}(0, 1), \text{seed} = 0)$; // evidence
 perturbation, fixed seed
- 4 $h_{\text{pr}} \leftarrow \text{DiffusionHop}_1(y, \text{seed} = 1)$; // prior perturbation, fixed
 evidence
- 5 $g_{\text{evidence}} \leftarrow |h_{\text{ev}} - h_0|$, $g_{\text{prior}} \leftarrow |h_{\text{pr}} - h_0|$;
- 6 $S_2 \leftarrow$ Eq. 8, upsampled to 1024×1024 ;
- 7 $(\hat{\sigma}_x, \hat{\sigma}_y, \hat{\theta}, \cdot) \leftarrow E_\theta(y)$;
- 8 $S_3 \leftarrow$ Eq. 9, broadcast to all pixels;
- 9 $f \leftarrow \phi(\hat{x}_1)$; // DINOv2 features of the first sample
- 10 $S_4 \leftarrow$ Eq. 10, broadcast to all pixels;
- 11 **return** $[S_1, S_2, S_3, S_4]$

3.4. Fusion head

The four raw signals, each independently rescaled to a comparable $[0, 1]$ -ish range before this stage, are stacked with one shallow image-luminance feature into a five-channel input $\mathbf{S}(p) = [S_1(p), S_2(p), S_3(p), S_4(p), L(p)]$ and passed through a small convolutional network f_θ : three 3×3 convolutional layers followed by two separate 1×1 heads, one producing a five-way per-pixel classification logit $f_\theta^{\text{cause}}(\mathbf{S})(p) \in \mathbb{R}^5$ over $\mathcal{C}$ (four causes plus a reliable class), the other a sigmoid-activated reliability score $f_\theta^{\text{rel}}(\mathbf{S})(p) \in [0, 1]$. Writing $y(p) \in \mathcal{C}$ for xSR-CausalBench’s ground-truth label at pixel p and Ω for the set of all labeled training pixels, both heads are trained jointly by minimizing

$$\mathcal{L}(\theta) = \frac{1}{|\Omega|} \sum_{p \in \Omega} \left[\underbrace{-\log \text{softmax}(f_\theta^{\text{cause}}(\mathbf{S})(p))_{y(p)}}_{\text{cause cross-entropy}} + \underbrace{\text{BCE}(f_\theta^{\text{rel}}(\mathbf{S})(p), \mathbb{I}[y(p) = \text{RELIABLE}])}_{\text{reliability binary cross-entropy}} \right], \quad (11)$$

with $\hat{c}(p)$ of Equation 3 recovered at inference as $\arg \max_c \text{softmax}(f_\theta^{\text{cause}}(\mathbf{S})(p))_c$.

This is the only trained component in the entire pipeline; the backbone, the kernel estimator, and the DINOv2 encoder are all frozen or independently pre-trained and contribute no term to Equation 11.

3.5. xSR-CausalBench

Training and evaluating the fusion head requires per-pixel labels for which cause dominates at each pixel, and because the true cause of a hallucination in a real photograph is unobservable, these labels must come from controlled synthetic construction rather than from annotation. Building on HalluGen’s diffusion-posterior-sampling controllable-hallucination idea [1], four procedures each take a held-out high-resolution source image and produce a low-resolution and high-resolution pair with a weak per-pixel label mask over five classes (the four causes plus reliable):

Ill-posedness-dominant degrades the entire image with the backbone’s standard, correctly specified degradation model at the full $16\times$ factor, labeling the whole image ill-posed by construction, since extreme downsampling alone is sufficient to destroy most of the null-space content regardless of anything else.

360 *Prior-reliance-dominant* degrades the image with the same standard model,
but first zeroes out the evidence in a rectangular sub-region covering 30% of the
image’s height and width (9% of its area), placed at a uniformly random valid
offset, beyond what the nominal degradation alone would remove; the exact
zeroed rectangle, not a fixed or scale-rounded approximation of it, is labeled
365 prior-reliance-dominant, and the remainder reliable.

Degradation-mismatch-dominant degrades the entire image with an out-of-
training-distribution kernel, anisotropic and rotated with a sigma range disjoint
from and above the standard pool, while content stays in-distribution, labeling
the whole image degradation-mismatch-dominant.

370 *Distribution-shift-dominant* blends a small patch, 64×64 pixels (0.4% of the
 1024×1024 image area), into an otherwise normal image before applying the
standard degradation model, then labels exactly the blended region distribution-
shift-dominant and the remainder reliable. The patch content is drawn from
a pool of natural images held disjoint from every procedure’s source images
375 (never from the same image being degraded), rather than from any exotic or
synthetically out-of-domain texture source; this is a proxy for distribution shift
using ordinary but contextually foreign content, not a claim that the patch
content is itself unusual in any absolute sense, and Section 5 treats the strength
of this proxy as an open question.

380 Each source image, drawn from a held-out validation split, is passed through
all four procedures independently and deterministically per index, so the bench-
mark is reproducible given a fixed seed. Two of the four procedures assign a label
from the physical operation applied (ill-posedness and degradation-mismatch,
both whole-image), while the other two assign a label from where evidence
385 was deliberately altered relative to the standard degradation (prior-reliance and
distribution-shift, both region-local); this means the same standard degradation
operator underlies the ill-posedness procedure’s whole-image ill-posed label and
the prior-reliance procedure’s remainder-region reliable label, with the two pro-
cedures differentiated only by whether the sub-region evidence was additionally
390 destroyed before that shared degradation was applied. This is a real construct-

validity limitation, not a subtlety hidden from the reader: the ground-truth label reflects which procedure generated the pixel, not an independently verified causal judgment about that pixel in isolation, and Section 5 returns to it directly. These four causes were not expected to separate perfectly; Section 5 reports the resulting entanglement directly. Algorithm 3 states the four-procedure dispatch precisely enough to reimplement without reading the source: given a fixed per-item seed, exactly one of the four procedures runs, deterministically, so re-generating the benchmark from the same (image list, seed) pair reproduces every pixel of every label mask exactly.

Algorithm 3: xSR-CausalBench sample construction

Input : HR source image x , OOD patch pool $\mathcal{P}$, procedure

$c \in \{\text{ill_posed}, \text{prior_reliance}, \text{mismatch}, \text{distribution_shift}\},$
seed s

Output: LR/HR pair (y, x') , per-pixel label map $L \in \{0, \dots, 4\}^{H \times W}$

```
1 rng ← Generator(s);
2 switch  $c$  do
3   case ill_posed do
4      $y \leftarrow \text{degrade\_standard}(x, 16, \text{rng}); x' \leftarrow x;$ 
5      $L \leftarrow \text{ILL\_POSED}$  everywhere;
6   end
7   case prior_reliance do
8      $(p_h, p_w) \leftarrow (0.3H, 0.3W); (y_0, x_0) \leftarrow \text{uniform random valid offset};$ 
9      $\tilde{x} \leftarrow x$  with region  $[y_0:y_0+p_h, x_0:x_0+p_w]$  set to 0;
10     $y \leftarrow \text{degrade\_standard}(\tilde{x}, 16, \text{rng}); x' \leftarrow x;$ 
11     $L \leftarrow \text{RELIABLE}$  everywhere, then PRIOR_RELIANCE in that
        region;
12  end
13  case mismatch do
14     $y \leftarrow \text{degrade\_mismatch}(x, 16, \text{rng});$  // OOD kernel:  $\sigma$  above
        standard range,  $\theta \neq 0$ 
15     $x' \leftarrow x; L \leftarrow \text{DEGRADATION\_MISMATCH}$  everywhere;
16  end
17  case distribution_shift do
18     $p \leftarrow \mathcal{P}[\text{image\_index mod } |\mathcal{P}|];$  // 64×64 patch, round-robin
19     $(y_0, x_0) \leftarrow \text{uniform random valid offset for } p;$ 
20     $\tilde{x} \leftarrow x$  with region  $[y_0:y_0+64, x_0:x_0+64]$  replaced by  $p;$ 
21     $y \leftarrow \text{degrade\_standard}(\tilde{x}, 16, \text{rng}); x' \leftarrow \tilde{x};$ 
22     $L \leftarrow \text{RELIABLE}$  everywhere, then DISTRIBUTION_SHIFT in
        that region;
23  end
24 end
25 return  $(y, x'), L$ 
```

400 4. Experiments

4.1. Setup

All experiments use the $16\times$ backbone described in Section 3.2, unmodified from its published checkpoint. The blind kernel estimator (Signal 3) was trained for 20 epochs on 500 images from the DIV2K training split [20], resized
405 to 256×256 patches, with the 50/50 standard/mismatch degradation mixture described in Section 3.3. The DINOv2 reference feature bank (Signal 4) was built once from 800 DIV2K training images (a superset of the kernel estimator’s 500), producing an 800-entry bank of 384-dimensional features. The fusion head was trained for 30 epochs on xSR-CausalBench instances built from 40
410 DIV2K validation images (160 training items across the four procedures, $K = 2$ stochastic backbone samples per item), with a held-out set of 20 further DIV2K validation images, disjoint from the training images, reserved for evaluation.

This scale is deliberately modest and is reported as such rather than disguised. The evaluation protocol below, including the expectation that the four
415 causes would not separate perfectly, was fixed before any results were collected; its target benchmark scale, 500 to 1000 source images per procedure, drawn from DIV2K validation together with Urban100 and Manga109 for structural and textural diversity, is one specific target this pilot falls short of by design, using a fraction of that scale so the full pipeline, including diffusion sampling
420 at every training item, could be validated end to end within the wall-clock budget available on a single consumer GPU. Section 5 treats closing that gap as a first-class limitation rather than a footnote.

4.2. Fidelity

Table 1 reports PSNR, SSIM, and LPIPS [21] between the super-resolved
425 output $\hat{x} \in [0, 1]^{3\times H\times W}$ and the true high-resolution image x , averaged over the same held-out evaluation items used throughout Section 4, for CHASR’s frozen diffusion backbone and, as a real fidelity comparison point, Real-ESRGAN [10] chained twice ($4\times$ then $4\times$, mirroring the diffusion backbone’s own two-hop

structure in Section 3.2) using its official public `RealESRGAN_x4plus` weights
 430 run on the identical inputs. PSNR is computed in the standard way from the
 mean squared error,

$$\text{PSNR}(\hat{x}, x) = 10 \log_{10} \left(\frac{1}{\text{MSE}(\hat{x}, x)} \right), \quad \text{MSE}(\hat{x}, x) = \frac{1}{3HW} \sum_{c,i,j} (\hat{x}_{c,i,j} - x_{c,i,j})^2, \quad (12)$$

with the numerator equal to 1 rather than the more familiar MAX^2 because
 both images are normalized to $[0, 1]$ before comparison. SSIM is the standard
 windowed luminance–contrast–structure product [21] and LPIPS is a learned
 435 perceptual distance in a pretrained network’s feature space rather than a closed-
 form pixel statistic; both are computed with their respective standard pub-
 lic implementations rather than re-derived here. Consistent with the framing
 throughout this paper, these are context metrics rather than the headline result:
 they are included because the special issue’s own evaluation protocol expects
 440 them, and because LPIPS specifically addresses this paper’s own stated concern
 that PSNR/SSIM under-penalize perceptually implausible content, not because
 a low pixel-fidelity score at $16\times$ from a diffusion prior is evidence of anything
 gone wrong. A model that reconstructs the exact original pixels of information
 the degradation destroyed would not be doing generative super-resolution; it
 445 would be doing something closer to lossless compression, which is not what a
 $16\times$ downsampling operator permits. The gap between these numbers and the
 higher PSNR/SSIM figures typically reported for diffusion super-resolution at
 $2\times$ to $4\times$ factors is exactly the phenomenon this paper is about, not a competing
 claim to be reconciled against it.

450 The comparison is genuinely mixed rather than one-sided, and reported
 here exactly as it came out. The diffusion backbone reaches a higher PSNR
 and a comparable SSIM than chained Real-ESRGAN, but a substantially worse
 (higher) LPIPS, 0.6730 versus 0.4977 — Real-ESRGAN’s output is judged per-
 ceptually closer to the true image despite scoring worse on both pixel-wise met-
 455 rics. This is consistent with, rather than a contradiction of, this paper’s central
 concern: a GAN trained with a supervised regression objective on synthetic

Table 1: Fidelity metrics on the held-out xSR-CausalBench items at $16\times$ (mean $\pm$ standard deviation across items), comparing CHASR’s frozen diffusion backbone against Real-ESRGAN chained to the same $16\times$ factor.

Metric	CHASR backbone	Real-ESRGAN (chained)
PSNR (dB)	19.16 ± 3.31	17.68 ± 3.28
SSIM	0.4194 ± 0.1714	0.4149 ± 0.2034
LPIPS (AlexNet)	0.6730 ± 0.1280	0.4977 ± 0.1349

degradations tends toward a more conservative, less hallucinated output at an extreme factor it was never designed for, while the diffusion prior fills the same null space with more confident, higher-frequency detail that raises pixel-wise fidelity in some places but perceptually diverges from the true image more often elsewhere. Neither number says which model is more useful for a given deployment; they say the two approaches trade off fidelity along different axes at $16\times$, which is exactly the kind of trade-off a pixel-level attribution method needs to explain rather than average away.

4.3. Attribution accuracy

Table 2 reports the fusion head’s five-way per-pixel cause classification accuracy and mean intersection-over-union against xSR-CausalBench’s synthetic ground truth on the held-out set. Writing $\hat{c}(p)$ for the predicted class of Equation 3 and $y(p)$ for xSR-CausalBench’s ground-truth label at pixel p , both over an evaluation image’s pixel set P , pixel accuracy is the fraction of pixels the predicted cause map recovers exactly,

$$\text{Acc} = \frac{1}{|P|} \sum_{p \in P} \mathbb{I}[\hat{c}(p) = y(p)], \quad (13)$$

and mean intersection-over-union averages, over the five-class taxonomy $\mathcal{C}$ of Equation 3, the overlap between the predicted region $\hat{P}_c = \{p : \hat{c}(p) = c\}$ and the true region $P_c = \{p : y(p) = c\}$ for each class,

$$\text{mIoU} = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \frac{|\hat{P}_c \cap P_c|}{|\hat{P}_c \cup P_c|}, \quad (14)$$

475 with the convention that $0/0 := 1$ for a class absent from both the prediction and
 the ground truth on a given image. Table 2 reports the per-image macro-average
 of Equation 14 unless stated otherwise; Table 4 reports the same quantity for
 every baseline compared there. xSR-CausalBench’s class prior is far from uni-
 form: two procedures (ill-posedness, mismatch) label an entire image with one
 480 class, while the other two (prior-reliance, distribution-shift) label only a small
 sub-region and leave the rest of that same image reliable, so reliable and the two
 whole-image classes dominate total pixel count and a uniform-random baseline
 (20%) understates how hard the task is to beat trivially. Table 2 therefore also
 reports the majority-class baseline, always predicting whichever class covers the
 485 most pixels in the held-out set (here, reliable, at 47.66%). The trained fusion
 head reaches 48.36% pixel accuracy, exceeding the majority-class baseline by
 0.7 percentage points and the uniform-random baseline by more than 28. This
 is a materially weaker result than comparing only against the uniform-random
 baseline suggests, and it is reported here rather than left for a reader to re-
 490 construct from the class distribution alone: at this training scale, most of the
 fusion head’s apparent accuracy is attributable to the class prior, not to gen-
 uinely learned per-pixel cause discrimination, and Section 5 treats this as the
 central open problem for scaling past this pilot, not a footnote to an otherwise
 strong number.

Table 2: Attribution accuracy on the held-out xSR-CausalBench split (5-way classification),
 against a uniform-random baseline and the empirical majority-class baseline (always predicting
 ‘reliable’, the most common ground-truth class in the held-out set).

Metric	Value
Pixel accuracy (fusion head)	48.36%
Pixel accuracy (majority-class baseline)	47.66%
Pixel accuracy (uniform-random baseline)	20.00%
Mean IoU, per-image macro-average	29.53%
Mean IoU, dataset-level accumulated confusion	19.50%

495 This narrow margin over the majority-class baseline has a direct, identifiable
 cause rather than being an unexplained shortfall. The fusion head was trained
 on only 160 items with unweighted cross-entropy, while two of the four pro-
 cedures (ill-posedness-dominant and degradation-mismatch-dominant) label an
 entire image with one class and the other two label only a small sub-region, so
 500 the two region-scale causes, prior-reliance and distribution-shift, are a small mi-
 nority of total training pixels with no class-balancing correction. An unweighted
 classifier trained under this imbalance is expected to converge toward something
 close to the majority-class rule, which is what Table 2 shows. Figure 2 makes
 this concrete with one held-out example per procedure: the predicted cause
 505 map for the prior-reliance example does recover the general location of the true
 suppressed region, but with substantial noise and incomplete coverage, exactly
 the signature of a model that has learned a real but weak signal on top of a
 strong class prior rather than either pure memorization of the prior or clean
 per-pixel discrimination. Separately, 40 training images is a small fraction of
 510 the design’s target scale, and the training loss curves (kernel estimator: 1.63
 dropping to a 0.87–0.93 plateau after one epoch; fusion head: 1.98 dropping
 to a 1.75–1.76 plateau after roughly eight epochs) both show real but modest
 learning consistent with a data-limited regime rather than a broken training
 signal. Section 5 treats class-balanced loss weighting, not just benchmark scale,
 515 as an immediate next step rather than optional future polish.

4.4. Disentanglement

Table 3 reports the Pearson correlation matrix among the four raw signal
 maps, flattened and pooled across all pixels of all 20 held-out images, to report
 the correlation directly rather than assume clean separation among the four
 520 causes. For two signal maps S_i, S_j flattened over the pooled pixel set P , with
 sample means $\bar{S}_i, \bar{S}_j$,

$$\rho(S_i, S_j) = \frac{\sum_{p \in P} (S_i(p) - \bar{S}_i)(S_j(p) - \bar{S}_j)}{\sqrt{\sum_{p \in P} (S_i(p) - \bar{S}_i)^2} \sqrt{\sum_{p \in P} (S_j(p) - \bar{S}_j)^2}}, \quad (15)$$

with $\rho \in [-1, 1]$ and $\rho(S_i, S_i) = 1$ on the diagonal by construction, which Table 3 reports for every pair $(i, j) \in \{1, 2, 3, 4\}^2$.

Table 3: Pearson correlation matrix among S1–S4 on the held-out evaluation set.

	S1 (null-space)	S2 (prior-reliance)	S3 (mismatch)	S4 (distribution shift)
S1	1.000	0.267	−0.060	−0.047
S2	0.267	1.000	−0.106	−0.123
S3	−0.060	−0.106	1.000	0.152
S4	−0.047	−0.123	0.152	1.000

Three of the six off-diagonal correlations are weak, at or below $|r| = 0.15$,
525 indicating the corresponding signal pairs behave close to independently across
the evaluation set. The exception is S1 and S2, correlated at $r = 0.267$: regions
with high null-space uncertainty also tend to show elevated prior reliance. This
is not a defect in the two signals; it is the expected consequence of what each
signal actually measures. A region with very little constraining evidence (high
530 S1) leaves the model with correspondingly less real signal to weigh against its
prior in the first place, which mechanically pushes prior reliance upward through
the same channel that produces the uncertainty. This coupling was flagged as
a risk before any data was collected; this experiment is the first evidence it is
real, and roughly the size the intuition predicted, rather than negligible or so
535 severe as to make the two signals redundant.

4.5. Ablation

Table 4 reports the ablation the evaluation protocol specifies: the trained
fusion head against each raw signal used alone and against an equal-weight com-
bination rule. Concretely, each single-signal baseline predicts that signal’s own
540 cause wherever it exceeds a threshold and reliable otherwise; the equal-weight
rule is the same idea applied jointly, predicting whichever of S1–S4 is largest at
a pixel if that maximum exceeds a shared threshold, reliable otherwise, since
it has no learned weights to combine the four signals any other way. Formally,

writing $S_{\max}(p) = \max_{i \in \{1, \dots, 4\}} S_i(p)$ for a candidate threshold rule and $\mathcal{D}_{\text{calib}}$
 545 for the fusion head’s own training set, every threshold τ_i reported in Table 4
 solves

$$\tau_i^* = \arg \max_{\tau \in \{0, 0.02, \dots, 1.0\}} \frac{1}{|\mathcal{D}_{\text{calib}}|} \sum_{p \in \mathcal{D}_{\text{calib}}} \mathbb{1}[\hat{c}_\tau(p) = y(p)], \quad \hat{c}_\tau(p) = \begin{cases} \arg \max_i S_i(p) & \text{if } S_{\max}(p) > \tau \\ \text{RELIABLE} & \text{otherwise,} \end{cases} \quad (16)$$

over the 51-point grid $\{0, 0.02, \dots, 1.0\}$, with the single-signal baselines a special
 case of Equation 16 restricted to one fixed i rather than the joint argmax.
 Every τ_i^* is fit on $\mathcal{D}_{\text{calib}}$ alone and then applied frozen to the held-out set for
 550 evaluation, so the comparison is never fit and evaluated on the same data.
 This calibration objective (raw accuracy under severe class imbalance) is itself
 a factor in what follows: a threshold that fires only on a rare cause typically
 costs more false positives against the dominant reliable class than it gains in
 true positives, so an accuracy-maximizing grid search is biased toward never
 555 firing regardless of how informative the underlying signal actually is, and the
 results below should be read with that bias in mind rather than as a clean
 measurement of signal informativeness alone. The fusion head’s own accuracy
 here (48.36%) reproduces Table 2’s number exactly; its mIoU (29.81%) differs
 slightly from Table 2’s 29.53% despite both coming from the same checkpoints
 560 and the same held-out set, a small discrepancy attributed to ordinary run-to-run
 floating-point non-determinism in fp16 diffusion sampling across two separate
 script executions rather than to any difference in method, and reported here
 rather than silently rounded away.

Three of the five baselines, equal-weight, S1 alone, and S2 alone, calibrate
 565 to a threshold of exactly 1.0, the upper boundary of the 51-point search grid
 rather than an interior optimum, meaning the grid search found that never firing
 (predicting reliable at every pixel) maximizes training-set accuracy at every grid
 point it tried. S4 alone calibrates to a genuine interior value (0.56), but S4’s own
 signal values on the held-out set never reach even that threshold, so it too never
 570 fires in practice; its accuracy and mIoU end up identical to the three threshold-

Table 4: Ablation: fusion head against every raw-signal and hand-combination baseline, on the held-out set (per-image macro-average).

Method	Pixel accuracy	Mean IoU
Fusion head (ours)	48.36%	29.81%
Equal-weight combination (threshold 1.00)	47.66%	23.83%
S1 (ill-posedness) alone (threshold 1.00)	47.66%	23.83%
S2 (prior-reliance) alone (threshold 1.00)	47.66%	23.83%
S3 (mismatch) alone (threshold 0.24)	46.52%	23.89%
S4 (distribution-shift) alone (threshold 0.56)	47.66%	23.83%

1.0 rules for that reason, not because its threshold is itself degenerate. All four therefore match the majority-class baseline in Table 2 exactly, because on this data they are the same rule. Only S3 finds a threshold (0.24) that actually fires, and it does slightly worse than the majority rule on accuracy (46.52%
575 vs. 47.66%, a 1.14-point absolute or 2.4% relative drop) while matching it on mIoU. The fusion head is the only method in this comparison that exceeds the majority-class rule on both metrics: by 0.70 points (1.5% relative) on accuracy, and by 5.98 points (25.1% relative) on mIoU, reporting both conventions here so neither metric’s margin is presented more favorably than the other. The larger
580 mIoU margin has an identifiable mechanical cause, not a different underlying effect: mIoU credits correctly identifying minority-class pixels (prior-reliance, distribution-shift) that every degenerate baseline scores zero intersection on by never predicting them at all, so accuracy, which is dominated by the numerically large reliable and whole-image classes, is far less sensitive to this difference than
585 mIoU is.

This comparison has two limits. First, it rules out five specific hand-crafted alternatives, not hand-crafted alternatives in general: a jointly-fit but still simple classifier on the same four raw signals, such as a per-pixel logistic regression, sits in the unexplored middle ground between a single grid-searched threshold
590 and the full convolutional fusion head, and this pilot does not test it. Second,

the margins above come from one run on 20 held-out images with no per-image variance or significance test reported, so they should be read as this pilot’s point estimate rather than as a result shown to be robust to resampling. Both are concrete, low-cost additions for the next iteration of this evaluation rather than
595 open-ended future work. With those two caveats attached, this paper reads the result as evidence the fusion head has learned something beyond the class prior among the alternatives actually tested here, not as evidence that no simpler method could ever match it: the fusion head has not been shown to be a strong classifier at this pilot’s scale, but it has been shown to be the only method among
600 those tested capable of attempting minority-class attribution at all, which is a real, if modest, form of value the raw signals and their naive combination do not provide on their own.




Figure 2: One held-out example per xSR-CausalBench procedure. LR input (nearest-neighbor upsampled for display), the backbone’s $16\times$ SR output, the ground-truth cause map, and the fusion head’s predicted cause map. Color key: ill-posed (red), prior-reliance (green), degradation-mismatch (yellow), distribution-shift (blue), reliable (gray). The prior-reliance row shows the predicted map recovering the approximate location of the true suppressed region with visible noise, consistent with the weak-but-real signal described above; the whole-image ill-posed and mismatch rows are largely uniform in both ground truth and prediction by construction.

5. Discussion and Limitations

5.1. *What this pilot validates, and what it does not*

605 The experiment reported here demonstrates that the CHASR architecture trains end to end on real images and produces a fusion head whose predictions are comfortably above uniform-random and, in the one qualitative example examined directly (Figure 2), spatially correlated with the true region rather than uniform noise; it also demonstrates that the disentanglement analysis specified
610 in advance produces an interpretable, honest result rather than an embarrassing surprise. It does not yet demonstrate that the fusion head has learned substantially more than the benchmark’s own class prior: the 0.7-point margin over the majority-class baseline (Table 2) is the single most important number in this paper to improve before any stronger claim is warranted, and framing
615 the pixel-accuracy result as state-of-the-art or as clearly beyond the class prior would misrepresent both the scale of the experiment and what the numbers can currently support. The training and evaluation sets used here, 40 and 20 images respectively, are roughly one twentieth of the design’s target scale of 500 to 1000 images per procedure. Scaling to that target, incorporating Urban100 and
620 Manga109 alongside DIV2K for structural and textural diversity as the design specifies, is the immediate and clearly defined next step, not a vague direction for future work.

One further evaluation component the design specifies was not run in this pilot and is named explicitly rather than left implicit: the downstream case
625 study correlating flagged high-mismatch or high-distribution-shift regions with actual optical character recognition failure on forensic-style text crops, which ties most directly to this special issue’s task-driven framing, was implemented as a reusable module but not exercised against real recognition failures in this pilot. This changes what can be claimed about CHASR’s practical value and
630 does not require a change to the architecture itself; it is evaluation work, not redesign work.

Section 4.5 supplies the ablation this same protocol calls for, comparing

the fusion head against each signal alone and against an equal-weight hand-combination rule. Its result sharpens, rather than resolves, the 0.70-point margin over the majority-class baseline flagged above (Table 2): none of the
635 five hand-crafted alternatives tested outperforms the trivial majority-class rule, while the fusion head is the only one that does, with a substantially larger margin on mean IoU than on accuracy. This paper reads that as evidence the fusion head has learned something beyond the class prior among the specific alterna-
640 tives tested, not as evidence that no simpler or differently-calibrated method ever could; Section 4.5 names the two most direct ways that stronger claim remains open (a jointly-fit simple classifier on the same signals, and per-image variance across a larger held-out set) rather than treating the ablation as closing the question.

Two measurement choices also warrant flagging directly. Signal 1’s variance
645 estimate is computed from $K = 2$ stochastic samples per pixel, a very small sample for a variance estimate; the qualitative pattern it produces is consistent with expectation, but the estimate itself is noisy at this K , and raising it is cheap (it costs additional backbone forward passes, not additional training) compared
650 to every other scaling step this section names. The mean IoU reported in Table 2 is a per-image macro-average, the mean of one mean_IoU computation per held-out image, rather than the dataset-level convention common in segmentation literature, which accumulates a single confusion matrix over the entire held-out set before dividing per class; the two are not numerically equivalent, and Table 2
655 reports the dataset-level number alongside it so this pilot’s mIoU is not read as more standardized than it is.

5.2. Factor entanglement

The $r = 0.267$ correlation between null-space uncertainty and prior reliance, reported without qualification in Section 4, means the four-way attribution this
660 paper proposes should not be read as claiming four cleanly independent causes that a classifier simply sorts a pixel into. It is closer to four correlated but distinguishable explanatory factors, in the same sense that a set of correlated predic-

tors in a regression can each still carry real, separable information even though they are not orthogonal. The practical consequence for a downstream user is that a pixel attributed primarily to prior reliance in a region that also shows high null-space uncertainty should be read as “the model is filling this in largely from its prior, and there was little evidence available for it to do otherwise,” a compound but coherent statement, rather than as two unrelated observations that happen to coincide. Whether this entanglement grows, shrinks, or restructures at the design’s full target scale is an open, empirical question this pilot’s sample size cannot answer with confidence, and is worth re-measuring once the benchmark is scaled rather than assumed to hold at the same magnitude.

5.3. *Synthetic-to-real gap*

Every ground-truth label used to train and evaluate the fusion head comes from xSR-CausalBench’s controlled synthetic construction, because the true cause of a hallucination in a genuine photograph is, by definition, not observable. HalluGen’s controllable-hallucination construction and the Hallucination Score’s severity judgments face the same constraint; no hallucination-attribution method plausibly avoids it. What can and should be strengthened is validating that performance on the synthetic benchmark transfers to genuinely uncontrolled degradations, using real low-resolution/high-resolution pairs from datasets such as RealSR or DRealSR where no synthetic cause label exists but severity correlation and qualitative inspection remain possible. This paper’s pilot evaluates on synthetic pairs only; the real-world generalization question remains open.

5.4. *Backbone dependency and forensic-use caveats*

CHASR’s four signals are computed against one specific frozen backbone, and both the calibration of Signal 3’s mismatch threshold and Signal 4’s reference feature bank are tied to that backbone’s specific training distribution. Swapping backbones, or fine-tuning the existing one, would require rebuilding the feature bank and likely re-tuning the mismatch calibration, though the

four-signal architecture and the fusion head training procedure would carry over unchanged. Readers considering CHASR or a similar attribution system for genuinely forensic use should also note that a calibrated reliability score reduces but does not eliminate the risk of over-trusting a hallucinated region: the fusion head’s reliability output is itself a learned prediction, trained on synthetic labels at pilot scale, and should be treated as an additional, imperfect signal to weigh alongside domain expertise rather than as a certification that a flagged-reliable pixel is ground truth.

6. Conclusion

Extreme super-resolution’s central practical problem is not that it hallucinates, since some degree of hallucination is an unavoidable consequence of the information the degradation destroys, but that existing systems give a user no way to distinguish an inherently unrecoverable region from a model that ignored evidence it should have used, from a mismatched operating condition, from genuinely unfamiliar content. This paper introduced a four-way causal taxonomy for exactly this distinction, the CHASR architecture that computes it as four signals fused by a small trained head sitting on top of an otherwise frozen backbone, and xSR-CausalBench, a synthetic benchmark that makes the four causes separately trainable and evaluable. A first end-to-end validation at real $16\times$ upscaling shows the architecture trains successfully and produces attribution accuracy well above uniform-random, though only narrowly above the benchmark’s own class prior at this pilot’s training scale, together with a disentanglement analysis that confirms, rather than merely speculates about, a factor-entanglement risk flagged before any data was collected, and an ablation that shows the trained fusion head, unlike every hand-crafted signal-threshold or equal-weight alternative tested, exceeds the benchmark’s own majority-class rule on both accuracy and mean IoU. The path from this pilot to a complete evaluation, class-balanced training at larger benchmark scale, and the downstream task-driven case study this special issue is centrally concerned with, is

concrete and already scoped, and is where this line of work goes next.

Data Availability Statement

The implementation, including the degradation pipelines, the four signal estimators, the fusion head, xSR-CausalBench construction code, and the training
725 and evaluation scripts used to produce every number reported in this paper, is available at https://github.com/vinhqdang/super_resolution. DIV2K [20] is a publicly available dataset used under its original terms.

Ethics Declaration

This study used only publicly available image data (DIV2K) and did not
730 involve human subjects, personal data, or any content requiring ethics board review.

Author Contributions (CRediT)

Q.D.: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review &
735 editing.

Declaration of Competing Interest

The author declares no competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

740 This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used artificial intelligence
745 tools to polish the manuscript text. After using these tools, the author reviewed
and edited the content as needed and takes full responsibility for the content of
the published article.

References

- [1] S. Kim, H. F. J. Tregidgo, C. Jin, M. Figini, D. C. Alexander, Hallugen:
750 Synthesizing realistic and controllable hallucinations for evaluating image
restoration, arXiv preprint arXiv:2512.03345 (2025).
- [2] W. Ren, R. Goyal, Z. Hu, T. T. Aumentado-Armstrong, I. Mohomed,
A. Levinshtein, Hallucination score: Towards mitigating hallucinations in
generative image super-resolution, arXiv preprint arXiv:2507.14367 (2025).
- 755 [3] S. Bell-Kligler, A. Shocher, M. Irani, Blind super-resolution kernel estima-
tion using an internal-gan, in: Advances in Neural Information Processing
Systems (NeurIPS), 2019.
- [4] J. Liang, G. Sun, K. Zhang, L. Van Gool, R. Timofte, Mutual affine network
for spatially variant kernel estimation in blind image super-resolution, in:
760 IEEE/CVF International Conference on Computer Vision (ICCV), 2021.
- [5] L. Wang, Y. Wang, X. Dong, Q. Xu, J. Yang, W. An, Y. Guo, Un-
supervised degradation representation learning for blind super-resolution,
in: IEEE/CVF Conference on Computer Vision and Pattern Recognition
(CVPR), 2021.
- 765 [6] J. Wang, Z. Yue, S. Zhou, K. C. Chan, C. C. Loy, Exploiting diffusion prior
for real-world image super-resolution, International Journal of Computer
Vision (2024).

- [7] R. Wu, T. Yang, L. Sun, Z. Zhang, S. Li, L. Zhang, SeeSR: Towards semantics-aware real-world image super-resolution, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024.
- [8] T. Yang, R. Wu, P. Ren, X. Xie, L. Zhang, Pixel-aware stable diffusion for realistic image super-resolution and personalized stylization, in: European Conference on Computer Vision (ECCV), 2024.
- [9] B. S. Kim, J. Kim, J. C. Ye, Chain-of-zoom: Extreme super-resolution via scale autoregression and preference alignment, in: Advances in Neural Information Processing Systems (NeurIPS), 2025, spotlight presentation. arXiv:2505.18600.
- [10] X. Wang, L. Xie, C. Dong, Y. Shan, Real-ESRGAN: Training real-world blind super-resolution with pure synthetic data, in: IEEE/CVF International Conference on Computer Vision Workshops (ICCVW), 2021.
- [11] R. Wang, Y.-Y. Chuang, Uncertainty-guided latent diffusion models for faithful super resolution, arXiv preprint arXiv:2608.25998 (2026).
- [12] J. Yin, J. Li, H. Xing, QUSR: Quality-aware and uncertainty-guided image super-resolution diffusion model, in: IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2026, arXiv:2603.09125.
- [13] K. N. H. Vo, A. Vazhentsev, A. Shelmanov, T. Baldwin, Y. Kementchedzhieva, Failing to see or failing to know? attributing errors in vision-language models, arXiv preprint arXiv:2607.04683 (2026).
- [14] Q.-V. Dang, N.-S.-A. Nguyen, T.-B.-D. Vo, M. N. Dinh, A. D. Le, H. V. Vu, P. L. Nguyen, Real-time hallucination correction in vision-language models using dynamic knowledge graph verification, Discover Artificial Intelligence (2026). doi:10.1007/s44163-026-01778-z.

- [15] S. K. Aithal, P. Maini, Z. C. Lipton, J. Z. Kolter, Understanding hallucinations in diffusion models through mode interpolation, in: Advances in Neural Information Processing Systems (NeurIPS), 2024.
- [16] Y. Wang, J. Yu, J. Zhang, Zero-shot image restoration using denoising diffusion null-space model, in: International Conference on Learning Representations (ICLR), 2023, oral presentation, notable-top-25%.
- [17] A. Sendjasni, M.-C. Larabi, Embedding similarity guided license plate super resolution, arXiv preprint arXiv:2501.01483 (2025).
- [18] K. S. Krishnan, K. S. Krishnan, HAT super-resolution and a PARSeq+CLIP4STR voting ensemble for extreme in-the-wild license plate recognition, arXiv preprint arXiv:2607.08896 (finalist entry, XLPSR Grand Challenge, ICIIP 2026) (2026).
- [19] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, et al., DINOv2: Learning robust visual features without supervision, Transactions on Machine Learning Research (2024).
- [20] E. Agustsson, R. Timofte, NTIRE 2017 challenge on single image super-resolution: Dataset and study, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 2017.
- [21] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, O. Wang, The unreasonable effectiveness of deep features as a perceptual metric, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2018.