

CHASR: Causal Hallucination Attribution for Extreme Super-Resolution

Quang-Vinh Dang

British University Vietnam, Hung Yen, Vietnam

Abstract

Extreme super-resolution (xSR, $8\times/16\times$ and beyond) is severely ill-posed: the diffusion priors that make plausible upscaling possible also make some hallucinated, plausible-but-false detail unavoidable. Existing work measures how much a super-resolved image hallucinates but not why, which limits its usefulness for forensic and task-driven applications needing to know whether a detail can be trusted. We introduce CHASR, a causal attribution module decomposing hallucination risk at the pixel level into four distinguishable causes: inherent ill-posedness, over-reliance on the generative prior, degradation-model mismatch, and distribution shift from the backbone’s training data. CHASR wraps a frozen chained diffusion backbone with four lightweight signal estimators and a small trained fusion head producing a per-pixel five-way cause classification plus a calibrated reliability score. We construct xSR-CausalBench, a synthetic benchmark injecting each cause under controlled conditions with weak per-pixel ground truth, extending an existing two-way controllable-hallucination construction to four causes. A first end-to-end validation at $16\times$ on real photographs (40 training, 20 held-out images) shows the fusion head reaches 48.4% five-way pixel attribution accuracy, above a uniform-random baseline (20%) but only 0.7 points above the benchmark’s empirical majority-class baseline (47.7%), a gap attributed to unweighted training on an imbalanced benchmark at pilot scale rather than the architecture. A disentanglement analysis confirms an an-

Email address: `vinh.dq4@buv.edu.vn` (Quang-Vinh Dang)

anticipated factor-entanglement risk: the null-space-uncertainty and prior-reliance signals correlate at $r = 0.267$, consistent with severe ill-posedness driving prior reliance. We report this pilot honestly, including where the numbers are weak. The central contribution is the causal taxonomy and attribution framework, not fidelity leaderboard performance.

Keywords: extreme super-resolution, hallucination attribution, diffusion models, causal analysis, null-space uncertainty, image forensics

1. Introduction

A security camera captures a license plate at a distance where the plate occupies perhaps a dozen pixels. A satellite image resolves a rooftop as a handful of gray values. A decades-old photograph, scanned at low resolution, is the only surviving record of a document. In each case, someone eventually asks a diffusion-based super-resolution system to recover detail that the sensor never captured, and in each case the system will comply: modern extreme super-resolution (xSR) methods, operating at $8\times$ or $16\times$ upscaling factors, now produce images that look convincingly sharp. The trouble is that “looks convincingly sharp” and “is actually correct” have become only loosely related properties. At these factors the information-theoretic content of the input is a small fraction of what the output claims to show, and the generative prior inside the model is doing most of the work of deciding what the missing pixels should be. When it decides wrong, the result is hallucination: detail that is plausible, consistent with the model’s training distribution, and false.

The severity of this problem scales directly with how the output will be used. A hallucinated license-plate character that happens to look legible is not a cosmetic defect if that character becomes evidence. A hallucinated texture on a satellite rooftop is not a rendering artifact if an analyst reports it as a structural feature. This is precisely the concern that motivates growing interest in extreme super-resolution’s forensic and task-driven applications, and it explains why the special issue this paper targets asks explicitly for evaluation metrics that go

beyond PSNR and SSIM: a fidelity score computed against a ground-truth image the deployment scenario does not have access to cannot tell an analyst whether
25 the ten pixels they are looking at right now can be trusted.

A handful of recent papers move past fidelity to measure hallucination directly. HalluGen [1] introduces a controllable benchmark that separates intrinsic hallucination (violations of data consistency) from extrinsic hallucination (wrong content filled into the model’s null space) and scores restorations against
30 this two-way taxonomy. The Hallucination Score [2] takes a complementary approach, using a multimodal large language model to produce a single scalar severity score that can double as a differentiable fine-tuning signal. Both are genuine advances, and both share a structural limitation for the use cases this paper is concerned with: they tell a user *how much* a region hallucinates, or
35 at best which of two broad categories the hallucination falls into, but not *why*. An analyst deciding whether to trust a specific license-plate character needs more than a severity number. They need to know whether the character is uncertain because the optics genuinely destroyed that information (in which case no model, however good, could have recovered it correctly), because the model
40 ignored usable evidence in favor of its own prior (a model defect, potentially fixable), because the model was applied to a degradation it was never trained to handle (an operating-conditions mismatch), or because the content itself is unlike anything the model has seen (an out-of-distribution failure). These are four distinguishable failure modes with four different practical responses, and
45 no existing method attributes hallucination to a specific one of them.

This paper introduces CHASR (Causal Hallucination Attribution for Super-Resolution), a module that sits on top of a frozen extreme super-resolution backbone and produces, for every pixel of the output, both a reliability score and a classification into one of four causes or a fifth “reliable” class. The design
50 commits to four causal signals that map directly onto the four failure modes above: null-space uncertainty, measured as the pixelwise variance across multiple degradation-consistent stochastic samples from the backbone; prior-reliance, measured as the relative sensitivity of the output to the low-resolution evidence

versus to the model’s random seed; degradation-model mismatch, measured by
55 re-estimating the input’s blur and noise parameters with an independent blind
kernel estimator and comparing them against the backbone’s assumed training
distribution; and distribution shift, measured as the nearest-neighbor distance
from the output’s frozen self-supervised features to a bank of the backbone’s
training-distribution features. A small fusion head, the only trained component
60 in the entire pipeline, learns to combine these four raw signals with a shallow
image feature into a calibrated per-pixel attribution.

Training and evaluating this fusion head requires per-pixel ground truth for
which cause produced a given hallucination, and no such labels exist for real
images because the true cause of a real hallucination is, definitionally, unobserv-
65 able. We address this by constructing xSR-CausalBench, a synthetic benchmark
that extends HalluGen’s controllable-hallucination construction from two causes
to four: four distinct procedures inject ill-posedness, prior-reliance, degradation-
mismatch, or distribution-shift dominance into an otherwise normal image un-
der experimentally controlled conditions, producing a weak per-pixel label mask
70 alongside each synthesized low-resolution and high-resolution pair.

The causes this taxonomy separates are not independent in practice, and
pretending otherwise would misrepresent what the framework can promise. A
region with extreme ill-posedness naturally invites more prior reliance, because
there is less real evidence available for the model to weigh against its prior in
75 the first place. Section 5 reports this correlation directly, measured from the
pipeline’s own output rather than assumed, because the alternative, quietly
hoping the four signals turn out to be independent and finding out otherwise at
review time, serves nobody.

This paper’s contribution is fourfold: a four-way causal taxonomy for ex-
80 treme super-resolution hallucination, extending the two-way intrinsic/extrinsic
split in the closest prior work; the CHASR architecture that instantiates the
taxonomy as four computable signals plus a learned fusion head, built on a pub-
licly reproducible diffusion backbone; xSR-CausalBench, a synthetic benchmark
that generates per-pixel weak ground truth for all four causes; and an honestly

85 reported pilot validation, at real $16\times$ upscaling on natural photographs, that demonstrates the pipeline is trainable and produces above-chance, structurally sensible attribution, while being explicit about what remains to scale before the numbers reported here should be read as a final answer rather than a first one.

2. Related Work

90 2.1. *Blind and diffusion-based super-resolution*

Classical blind super-resolution treats the degradation kernel as an unknown to be estimated jointly with, or prior to, the restoration itself. KernelGAN [3] estimates a per-image blur kernel using an internal GAN trained purely on patch recurrence within the low-resolution image, requiring no external training data. 95 MANet [4] instead learns a mutual affine network that estimates spatially variant kernels directly, relaxing the common assumption of a single global kernel per image. DASR [5] sidesteps explicit kernel parameterization entirely, learning an unsupervised degradation representation via contrastive learning and conditioning the restoration network on it directly. This lineage of blind kernel 100 estimation is a component CHASR reuses directly: Signal 3 (degradation-model mismatch) is computed by an independent blind estimator trained on both the backbone’s assumed degradation pool and an out-of-distribution pool, then compared against the backbone’s training assumptions, rather than proposed here as a new estimation method.

105 Diffusion models have displaced GAN-based approaches as the dominant paradigm for photorealistic super-resolution at moderate factors. StableSR [6] adapts a pretrained text-to-image diffusion prior for real-world restoration through a time-aware encoder and a controllable feature-wrapping module. SeeSR [7] conditions the diffusion process on semantic tags extracted from the degraded 110 input, improving fidelity to the true image content rather than merely to a plausible one. PASD [8] introduces pixel-aware cross-attention so that the diffusion process attends explicitly to which regions the input actually constrains. Chain-of-Zoom [9] pushes this diffusion-prior paradigm to substantially more extreme

factors by reformulating extreme super-resolution as an autoregressive sequence
of moderate zoom-in steps guided by vision-language model text prompts, each
115 step handled by an off-the-shelf $4\times$ diffusion model rather than a single model
retrained for the full factor. CHASR’s backbone follows this same chained-hop
principle for reaching $16\times$ from a $4\times$ checkpoint, though for reproducibility on
modest hardware it wraps the publicly downloadable `stable-diffusion-x4-u`
120 `pscaler` checkpoint directly rather than a VLM-guided variant, trading some of
Chain-of-Zoom’s content-awareness for a fully open, easily reproduced pipeline.
Real-ESRGAN [10] remains the standard GAN-based reference point for real-
world super-resolution trained entirely on synthetic degradations, and Section 4
reports it as a real fidelity baseline, chained to the same $16\times$ factor as CHASR’s
125 own backbone, rather than only citing it here. Two concurrent lines of work cou-
ple diffusion-based super-resolution directly to an uncertainty signal rather than
treating uncertainty as an afterthought: UGDiff [11] estimates per-latent-feature
reconstruction uncertainty and uses it to steer high-frequency detail restoration
toward high-uncertainty regions while leaving confident regions untouched, and
130 QUSR [12] adapts injected diffusion noise strength to local uncertainty along-
side a multimodal-LLM-based quality prior, under unknown, spatially varying
degradation. Both use uncertainty as a single scalar steering signal for where to
generate more aggressively; CHASR’s Signal 1 measures a structurally similar
null-space quantity but treats it as one of four distinguishable causes to attribute
135 after generation rather than a knob to steer generation with, a diagnostic rather
than a generative use of the same underlying idea.

2.2. Measuring hallucination in generative restoration

HalluGen [1] is the closest prior work to CHASR’s evaluation methodology.
It separates hallucination into intrinsic (violating data consistency with the
140 observed degradation) and extrinsic (wrong content filled into the null space
where data consistency alone does not constrain the answer) categories, con-
structs a synthetic controllable-hallucination benchmark via diffusion posterior
sampling, and proposes a patch-level severity metric. Its own benchmark is built

from low-field MRI restoration rather than natural photographs, so CHASR’s
 145 xSR-CausalBench (Section 3.4) is a cross-domain methodological descendant
 rather than a direct reuse: it ports the controllable-injection idea from medi-
 cal to natural-image extreme super-resolution, extending it from two categories
 to four and from a severity score to a per-pixel causal label. The Halluci-
 nation Score [2] takes a different but complementary route: rather than de-
 150 composing hallucination by cause, it uses a multimodal large language model
 to judge two properties, implausibility under any degradation and semantic
 deviation from what the input evidence supports, and combines them into a
 scalar that doubles as a reward signal for fine-tuning. Both works measure
 severity or coarse category; neither attributes a specific hallucinated region
 155 to one of several distinguishable causes, which is the gap this paper targets
 directly. The same detection-versus-attribution gap has emerged concurrently
 outside super-resolution: Vo et al. [13] decompose vision-language model failures
 on knowledge-intensive visual question answering into recognition-stage versus
 knowledge-stage causes rather than a single failure score, enabling different fixes
 160 (image repair, entity support, question rewriting) for different causes, the same
 why-before-what argument this paper makes for extreme super-resolution. In a
 related line of the present author’s own work, addressing a structurally similar
 problem in a different modality, Dang et al. [14] intercept and correct halluci-
 nated tokens during vision-language model generation via a dynamically con-
 165 structed, externally-verified knowledge graph rather than a post-hoc severity
 score, underscoring that mechanism-aware, causal intervention is an emerging
 theme across generative-AI hallucination research generally, not a concern spe-
 cific to image restoration.

At a more general level, work on diffusion models outside restoration has
 170 begun explaining hallucination as an out-of-distribution phenomenon: Aithal et
 al. [15] show that diffusion models can hallucinate through mode interpolation,
 generating samples in regions between distinct modes of the training distribu-
 tion that the model never actually observed, which is a mechanistic account
 of why unfamiliar content produces confident-looking but wrong output. Sig-

nal 4 in this paper’s method operationalizes the restoration-specific analogue of that same intuition, measuring an output’s feature-space distance from the backbone’s training distribution directly, rather than proposing a new theory of why distribution shift causes hallucination.

Range-null space decomposition, formalized for diffusion-based zero-shot inverse problems by DDNM [16], provides the mathematical grounding for CHASR’s Signal 1. DDNM shows that for a known linear degradation operator, a diffusion sampling trajectory can be projected at every step so that the range-space component exactly matches the observation while the diffusion prior is left free to fill the null space, and demonstrates this on a range of restoration tasks under the assumption that the degradation operator is known in advance. CHASR uses the same range-null projection, applied to the average-pooling operator that governs its synthetic degradation pipeline, to guarantee that every stochastic sample the backbone produces is exactly degradation-consistent, which is what makes the sample-to-sample variance in Signal 1 attributable to null-space uncertainty specifically rather than to a data-consistency violation. Where DDNM assumes a known degradation and does not address attribution, CHASR assumes a partially unknown degradation (estimated blindly for Signal 3) and uses the resulting variance as one input to a diagnostic rather than as the whole restoration method.

2.3. Task-driven super-resolution and this special issue

The special issue this paper targets is organized in part around task-driven reliability, and closely related work on embedding-similarity-guided license-plate super-resolution [17] exemplifies the concern directly: a super-resolved license plate is only useful if the specific characters it renders are correct, not merely if the overall image looks sharp. The concurrent XLPSR Grand Challenge at ICIP 2026, whose finalist entries include ensemble approaches combining super-resolution backbones with scene-text recognition voting [18], reflects the same practical stakes at competition scale: entries are scored on downstream recognition accuracy, not perceptual quality. CHASR’s downstream case study,

205 correlating flagged high-mismatch or high-distribution-shift regions with actual optical character recognition failure on text-patch crops rather than treating task-driven evaluation as secondary to fidelity, is implemented as a reusable module (Section 5) but not yet exercised against real recognition failures in this pilot.

210 3. Method

CHASR is a diagnostic module, not a new super-resolution architecture. It wraps a frozen backbone and adds four causal signal estimators plus a small trained fusion head, so that the entire training cost of the system is confined to the fusion head while the backbone’s own super-resolution quality is left 215 exactly as published. Figure 1 shows the full pipeline: the frozen backbone (Section 3.1) produces the super-resolved output and the raw material each signal estimator (Section 3.2) needs, and the only trained component, the fusion head (Section 3.3), combines the four resulting signal maps into a per-pixel cause classification and reliability score.

220 3.1. Frozen backbone

The backbone reaches $16\times$ from 64×64 low-resolution patches through two chained $4\times$ diffusion hops, using the publicly downloadable checkpoint **stable-diffusion-x4-upscaler** (hosted on the Hugging Face Hub under the **stabilityai** organization) at each hop for full reproducibility on a single con- 225 sumer GPU. The first hop upscales the full 64×64 patch to 256×256 . The second hop is applied per-tile, splitting the 256×256 intermediate into four 128×128 tiles, each independently upscaled to 512×512 and stitched back into the final 1024×1024 output, which keeps peak memory within an 8 GB budget at the cost of a faint seam artifact at tile boundaries that does not affect the causal signals computed here. After every hop, the output is projected via the DDNM range- 230 null decomposition [16] against the true low-resolution observation at that hop’s scale, using the closed-form projection for an average-pooling degradation operator: within every scale-by-scale block, the block is replaced by itself minus its




Figure 1: CHASR pipeline. Blue: frozen backbone. Green: the four independent signal estimators, each reading the backbone’s output and/or intermediate state (Section 3.2). Orange: the fusion head, the only trained component, which stacks the four signal maps with a shallow luminance feature and predicts a per-pixel cause and reliability.

own mean plus the corresponding observed pixel. This guarantees every sample
 235 the backbone returns is exactly consistent with the input, at every hop, which
 is what makes the variance across independently sampled outputs attributable
 specifically to null-space content rather than to the model occasionally violating
 the observation. Algorithm 1 states this procedure precisely.

3.2. Four causal signals

240 **Signal 1: null-space uncertainty.** The backbone is sampled K times with
 different random seeds, each sample independently DDNM-projected as above.
 The pixelwise, channel-averaged unbiased variance across the K samples is

$$S_1(p) = \text{Var}_k[\hat{x}_k(p)]. \quad (1)$$

Algorithm 1: Chained, DDNM-projected 16 $\times$ sampling

Input : Low-resolution patch $y \in \mathbb{R}^{64 \times 64 \times 3}$, sample count K

Output: K degradation-consistent 16 $\times$ super-resolved samples

$$\{\hat{x}_1, \dots, \hat{x}_K\}$$

```
1 for  $k \leftarrow 1$  to  $K$  do
2    $z_1 \leftarrow \text{DiffusionHop}_1(y, \text{seed} = k)$  ;           //  $64 \times 64 \rightarrow 256 \times 256$ 
3    $z_1 \leftarrow \text{DDNM\_project}(z_1, y, \text{scale} = 4)$  ; // exact consistency at
   hop 1's scale
4   for  $\text{tile} \in \text{split}(z_1, 4)$  do
5      $t \leftarrow \text{DiffusionHop}_2(\text{tile}, \text{seed} = k)$  ;    //  $128 \times 128 \rightarrow 512 \times 512$ 
6      $t \leftarrow \text{DDNM\_project}(t, \text{tile}, \text{scale} = 4)$ ;
7     stitch( $t$ ) into  $z_2 \in \mathbb{R}^{1024 \times 1024 \times 3}$ ;
8   end
9    $\hat{x}_k \leftarrow \text{DDNM\_project}(z_2, y, \text{scale} = 16)$  ; // final consistency
   check against  $y$  directly
10 end
11 return  $\{\hat{x}_1, \dots, \hat{x}_K\}$ 
```

Because data consistency is fixed by construction, the samples can only differ in their null-space content, so high variance at a pixel means the ill-posedness at that pixel is large enough that the model genuinely could not have determined a unique answer from the evidence, independent of whether the model itself is well-calibrated.

Signal 2: generative-prior over-reliance. This signal asks a different question: even where evidence is available, did the model use it? For the frozen diffusion backbone, whose stochasticity is a discrete random seed rather than a continuous latent, this is estimated by finite differences at the first (256 $\times$ 256) hop: the raw, pre-projection output's sensitivity to a small perturbation of the low-resolution evidence (g_{evidence}) is compared against its sensitivity to a full

reseed at fixed evidence (g_{prior}),

$$S_2(p) = \frac{g_{\text{prior}}(p)}{g_{\text{prior}}(p) + g_{\text{evidence}}(p) + \varepsilon}. \quad (2)$$

Both terms are compared as raw output-magnitude changes on the same units, rather than as normalized derivatives, because normalizing only the continuous perturbation by its step size while leaving the discrete reseed unnormalized would bias the ratio toward zero regardless of the model’s actual behavior. A general finite-difference formulation for a continuously perturbable forward function and latent is also implemented as a standalone, independently tested utility for future backbones with a true continuous latent, but is not used in the real pipeline for exactly this reason.

Signal 3: degradation-model mismatch. An independent, lightweight convolutional network is trained to regress a blur kernel’s spatial parameters (anisotropic sigma $(\hat{\sigma}_x, \hat{\sigma}_y)$, rotation angle $\hat{\theta}$) and additive noise level directly from a 64×64 patch, trained on a 50/50 mixture of the backbone’s assumed in-distribution degradation pool (upper sigma bound σ_{max} , zero rotation) and an out-of-distribution pool with a disjoint, larger sigma range and non-zero rotation. At inference, this estimator re-examines the actual input patch and scores how far the estimate falls outside the assumed pool,

$$S_3 = \min\left(1, \frac{\max(0, \max(\hat{\sigma}_x, \hat{\sigma}_y) - \sigma_{\text{max}}) + |\hat{\theta}|}{d_{\text{max}}}\right), \quad (3)$$

normalized by the out-of-distribution pool’s own worst-case raw distance d_{max} so the signal saturates at 1.0 only at genuinely extreme mismatch rather than prematurely. S_3 is a single scalar per input patch, broadcast uniformly to every pixel, since the degradation kernel is assumed spatially constant within one patch. A high value here means the backbone is being asked to restore an image degraded in a way it was never exposed to during training, independent of whether the content itself is familiar.

Signal 4: distribution shift. A frozen, self-supervised DINOv2 encoder [19] extracts a feature vector f from the super-resolved output, and its k -nearest-neighbor distance $d_{\text{kNN}}(f, \mathcal{B})$ to a reference bank $\mathcal{B}$ of features extracted once

from the backbone’s training-distribution images is computed. Because a sparse reference bank in a high-dimensional feature space suffers a curse-of-dimensionality effect at $k > 1$, where a near-duplicate query’s second through k -th neighbors are effectively generic points at a roughly constant distance that washes out fine-grained separation, the raw distance is rescaled by the bank’s own internal
285 typical k -nearest-neighbor distance $s(\mathcal{B})$ among its own members,

$$S_4 = \frac{d_{\text{kNN}}(f, \mathcal{B})}{d_{\text{kNN}}(f, \mathcal{B}) + s(\mathcal{B})}, \quad (4)$$

producing a self-calibrated score in $[0, 1]$ without assuming a fixed feature-space magnitude in advance. Like S_3 , S_4 is computed once per output image and broadcast to every pixel.

290 Algorithm 2 collects Equations 1–4 in one place, making explicit which quantities each signal shares (all four read from the same K -sample set of Algorithm 1, and S_2 additionally requires two extra single-hop calls beyond it) and which are independently pretrained rather than computed fresh per pixel (S_3 ’s kernel estimator E_θ and S_4 ’s DINOv2 encoder ϕ are both frozen, queried here, not trained here).
295

3.3. Fusion head

The four raw signals, each independently rescaled to a comparable $[0, 1]$ -ish range before this stage, are stacked with one shallow image-luminance feature into a five-channel input and passed through a small convolutional network:
300 three 3×3 convolutional layers followed by two separate 1×1 heads, one producing a five-way per-pixel classification logit (four causes plus a reliable class) trained with cross-entropy against the synthetic benchmark’s weak labels, the other producing a sigmoid-activated reliability score trained with binary cross-entropy against whether the ground-truth label at that pixel is the reliable
305 class. This is the only trained component in the entire pipeline; the backbone, the kernel estimator, and the DINOv2 encoder are all frozen or independently pretrained.

Algorithm 2: Four-signal computation

Input : LR patch y , samples $\{\hat{x}_1, \dots, \hat{x}_K\}$ from Algorithm 1, kernel estimator E_θ , DINOv2 encoder ϕ , reference bank $\mathcal{B}$

Output: Signal stack $[S_1, S_2, S_3, S_4] \in [0, 1]^{4 \times 1024 \times 1024}$

- 1 $S_1 \leftarrow \text{Var}_k[\hat{x}_k]$ via Eq. 1, averaged over channels, normalized to $[0, 1]$;
- 2 $h_0 \leftarrow \text{DiffusionHop}_1(y, \text{seed} = 0)$;
- 3 $h_{\text{ev}} \leftarrow \text{DiffusionHop}_1(y + \epsilon \cdot \mathcal{N}(0, 1), \text{seed} = 0)$; // evidence
 perturbation, fixed seed
- 4 $h_{\text{pr}} \leftarrow \text{DiffusionHop}_1(y, \text{seed} = 1)$; // prior perturbation, fixed
 evidence
- 5 $g_{\text{evidence}} \leftarrow |h_{\text{ev}} - h_0|$, $g_{\text{prior}} \leftarrow |h_{\text{pr}} - h_0|$;
- 6 $S_2 \leftarrow$ Eq. 2, upsampled to 1024×1024 ;
- 7 $(\hat{\sigma}_x, \hat{\sigma}_y, \hat{\theta}, \cdot) \leftarrow E_\theta(y)$;
- 8 $S_3 \leftarrow$ Eq. 3, broadcast to all pixels;
- 9 $f \leftarrow \phi(\hat{x}_1)$; // DINOv2 features of the first sample
- 10 $S_4 \leftarrow$ Eq. 4, broadcast to all pixels;
- 11 **return** $[S_1, S_2, S_3, S_4]$

3.4. *xSR-CausalBench*

Training and evaluating the fusion head requires per-pixel labels for which
310 cause dominates at each pixel, and because the true cause of a hallucination
in a real photograph is unobservable, these labels must come from controlled
synthetic construction rather than from annotation. Building on HalluGen’s
diffusion-posterior-sampling controllable-hallucination idea [1], four procedures
each take a held-out high-resolution source image and produce a low-resolution
315 and high-resolution pair with a weak per-pixel label mask over five classes (the
four causes plus reliable):

Ill-posedness-dominant degrades the entire image with the backbone’s stan-
dard, correctly specified degradation model at the full $16\times$ factor, labeling the
whole image ill-posed by construction, since extreme downsampling alone is

320 sufficient to destroy most of the null-space content regardless of anything else.

Prior-reliance-dominant degrades the image with the same standard model, but first zeroes out the evidence in a rectangular sub-region covering 30% of the image’s height and width (9% of its area), placed at a uniformly random valid offset, beyond what the nominal degradation alone would remove; the exact
325 zeroed rectangle, not a fixed or scale-rounded approximation of it, is labeled prior-reliance-dominant, and the remainder reliable.

Degradation-mismatch-dominant degrades the entire image with an out-of-training-distribution kernel, anisotropic and rotated with a sigma range disjoint from and above the standard pool, while content stays in-distribution, labeling
330 the whole image degradation-mismatch-dominant.

Distribution-shift-dominant blends a small patch, 64×64 pixels (0.4% of the 1024×1024 image area), into an otherwise normal image before applying the standard degradation model, then labels exactly the blended region distribution-shift-dominant and the remainder reliable. The patch content is drawn from
335 a pool of natural images held disjoint from every procedure’s source images (never from the same image being degraded), rather than from any exotic or synthetically out-of-domain texture source; this is a proxy for distribution shift using ordinary but contextually foreign content, not a claim that the patch content is itself unusual in any absolute sense, and Section 5 treats the strength
340 of this proxy as an open question.

Each source image, drawn from a held-out validation split, is passed through all four procedures independently and deterministically per index, so the benchmark is reproducible given a fixed seed. Two of the four procedures assign a label from the physical operation applied (ill-posedness and degradation-mismatch,
345 both whole-image), while the other two assign a label from where evidence was deliberately altered relative to the standard degradation (prior-reliance and distribution-shift, both region-local); this means the same standard degradation operator underlies the ill-posedness procedure’s whole-image ill-posed label and the prior-reliance procedure’s remainder-region reliable label, with the two pro-
350 cedures differentiated only by whether the sub-region evidence was additionally

destroyed before that shared degradation was applied. This is a real construct-
validity limitation, not a subtlety hidden from the reader: the ground-truth
label reflects which procedure generated the pixel, not an independently veri-
fied causal judgment about that pixel in isolation, and Section 5 returns to it
355 directly. These four causes were not expected to separate perfectly; Section 5 re-
ports the resulting entanglement directly. Algorithm 3 states the four-procedure
dispatch precisely enough to reimplement without reading the source: given a
fixed per-item seed, exactly one of the four procedures runs, deterministically,
so re-generating the benchmark from the same (image list, seed) pair reproduces
360 every pixel of every label mask exactly.

Algorithm 3: xSR-CausalBench sample construction

Input : HR source image x , OOD patch pool $\mathcal{P}$, procedure

$c \in \{\text{ill_posed}, \text{prior_reliance}, \text{mismatch}, \text{distribution_shift}\},$
seed s

Output: LR/HR pair (y, x') , per-pixel label map $L \in \{0, \dots, 4\}^{H \times W}$

```
1 rng ← Generator(s);
2 switch  $c$  do
3   case ill_posed do
4      $y \leftarrow \text{degrade\_standard}(x, 16, \text{rng}); x' \leftarrow x;$ 
5      $L \leftarrow \text{ILL\_POSED}$  everywhere;
6   end
7   case prior_reliance do
8      $(p_h, p_w) \leftarrow (0.3H, 0.3W); (y_0, x_0) \leftarrow \text{uniform random valid offset};$ 
9      $\tilde{x} \leftarrow x$  with region  $[y_0:y_0+p_h, x_0:x_0+p_w]$  set to 0;
10     $y \leftarrow \text{degrade\_standard}(\tilde{x}, 16, \text{rng}); x' \leftarrow x;$ 
11     $L \leftarrow \text{RELIABLE}$  everywhere, then PRIOR_RELIANCE in that
        region;
12  end
13  case mismatch do
14     $y \leftarrow \text{degrade\_mismatch}(x, 16, \text{rng});$  // OOD kernel:  $\sigma$  above
        standard range,  $\theta \neq 0$ 
15     $x' \leftarrow x; L \leftarrow \text{DEGRADATION\_MISMATCH}$  everywhere;
16  end
17  case distribution_shift do
18     $p \leftarrow \mathcal{P}[\text{image\_index mod } |\mathcal{P}|];$  // 64×64 patch, round-robin
19     $(y_0, x_0) \leftarrow \text{uniform random valid offset for } p;$ 
20     $\tilde{x} \leftarrow x$  with region  $[y_0:y_0+64, x_0:x_0+64]$  replaced by  $p;$ 
21     $y \leftarrow \text{degrade\_standard}(\tilde{x}, 16, \text{rng}); x' \leftarrow \tilde{x};$ 
22     $L \leftarrow \text{RELIABLE}$  everywhere, then DISTRIBUTION_SHIFT in
        that region;
23  end
24 end
25 return  $(y, x'), L$ 
```

4. Experiments

4.1. Setup

All experiments use the $16\times$ backbone described in Section 3.1, unmodified from its published checkpoint. The blind kernel estimator (Signal 3) was trained for 20 epochs on 500 images from the DIV2K training split [20], resized to 256×256 patches, with the 50/50 standard/mismatch degradation mixture described in Section 3.2. The DINOv2 reference feature bank (Signal 4) was built once from 800 DIV2K training images (a superset of the kernel estimator’s 500), producing an 800-entry bank of 384-dimensional features. The fusion head was trained for 30 epochs on xSR-CausalBench instances built from 40 DIV2K validation images (160 training items across the four procedures, $K = 2$ stochastic backbone samples per item), with a held-out set of 20 further DIV2K validation images, disjoint from the training images, reserved for evaluation.

This scale is deliberately modest and is reported as such rather than disguised. The evaluation protocol below, including the expectation that the four causes would not separate perfectly, was fixed before any results were collected; its target benchmark scale, 500 to 1000 source images per procedure, drawn from DIV2K validation together with Urban100 and Manga109 for structural and textural diversity, is one specific target this pilot falls short of by design, using a fraction of that scale so the full pipeline, including diffusion sampling at every training item, could be validated end to end within the wall-clock budget available on a single consumer GPU. Section 5 treats closing that gap as a first-class limitation rather than a footnote.

4.2. Fidelity

Table 1 reports PSNR, SSIM, and LPIPS [21] between the super-resolved output and the true high-resolution image, averaged over the same held-out evaluation items used throughout Section 4, for CHASR’s frozen diffusion backbone and, as a real fidelity comparison point, Real-ESRGAN [10] chained twice ($4\times$ then $4\times$, mirroring the diffusion backbone’s own two-hop structure in Section 3.1) using its official public `RealESRGAN_x4plus` weights run on the identical

inputs. Consistent with the framing throughout this paper, these are context metrics rather than the headline result: they are included because the special issue’s own evaluation protocol expects them, and because LPIPS specifically addresses this paper’s own stated concern that PSNR/SSIM under-penalize perceptually implausible content, not because a low pixel-fidelity score at 16× from a diffusion prior is evidence of anything gone wrong. A model that reconstructs the exact original pixels of information the degradation destroyed would not be doing generative super-resolution; it would be doing something closer to lossless compression, which is not what a 16× downsampling operator permits. The gap between these numbers and the higher PSNR/SSIM figures typically reported for diffusion super-resolution at 2× to 4× factors is exactly the phenomenon this paper is about, not a competing claim to be reconciled against it.

Table 1: Fidelity metrics on the held-out xSR-CausalBench items at 16× (mean $\pm$ standard deviation across items), comparing CHASR’s frozen diffusion backbone against Real-ESRGAN chained to the same 16× factor.

Metric	CHASR backbone	Real-ESRGAN (chained)
PSNR (dB)	19.16 $\pm$ 3.31	17.68 $\pm$ 3.28
SSIM	0.4194 $\pm$ 0.1714	0.4149 $\pm$ 0.2034
LPIPS (AlexNet)	0.6730 $\pm$ 0.1280	0.4977 $\pm$ 0.1349

The comparison is genuinely mixed rather than one-sided, and reported here exactly as it came out. The diffusion backbone reaches a higher PSNR and a comparable SSIM than chained Real-ESRGAN, but a substantially worse (higher) LPIPS, 0.6730 versus 0.4977 — Real-ESRGAN’s output is judged perceptually closer to the true image despite scoring worse on both pixel-wise metrics. This is consistent with, rather than a contradiction of, this paper’s central concern: a GAN trained with a supervised regression objective on synthetic degradations tends toward a more conservative, less hallucinated output at an extreme factor it was never designed for, while the diffusion prior fills the same null space with more confident, higher-frequency detail that raises pixel-wise

fidelity in some places but perceptually diverges from the true image more often elsewhere. Neither number says which model is more useful for a given deployment; they say the two approaches trade off fidelity along different axes at $16\times$, which is exactly the kind of trade-off a pixel-level attribution method needs to explain rather than average away.

4.3. Attribution accuracy

Table 2 reports the fusion head’s five-way per-pixel cause classification accuracy and mean intersection-over-union against xSR-CausalBench’s synthetic ground truth on the held-out set. xSR-CausalBench’s class prior is far from uniform: two procedures (ill-posedness, mismatch) label an entire image with one class, while the other two (prior-reliance, distribution-shift) label only a small sub-region and leave the rest of that same image reliable, so reliable and the two whole-image classes dominate total pixel count and a uniform-random baseline (20%) understates how hard the task is to beat trivially. Table 2 therefore also reports the majority-class baseline, always predicting whichever class covers the most pixels in the held-out set (here, reliable, at 47.66%). The trained fusion head reaches 48.36% pixel accuracy, exceeding the majority-class baseline by 0.7 percentage points and the uniform-random baseline by more than 28. This is a materially weaker result than comparing only against the uniform-random baseline suggests, and it is reported here rather than left for a reader to reconstruct from the class distribution alone: at this training scale, most of the fusion head’s apparent accuracy is attributable to the class prior, not to genuinely learned per-pixel cause discrimination, and Section 5 treats this as the central open problem for scaling past this pilot, not a footnote to an otherwise strong number.

This narrow margin over the majority-class baseline has a direct, identifiable cause rather than being an unexplained shortfall. The fusion head was trained on only 160 items with unweighted cross-entropy, while two of the four procedures (ill-posedness-dominant and degradation-mismatch-dominant) label an entire image with one class and the other two label only a small sub-region, so

Table 2: Attribution accuracy on the held-out xSR-CausalBench split (5-way classification), against a uniform-random baseline and the empirical majority-class baseline (always predicting ‘reliable’, the most common ground-truth class in the held-out set).

Metric	Value
Pixel accuracy (fusion head)	48.36%
Pixel accuracy (majority-class baseline)	47.66%
Pixel accuracy (uniform-random baseline)	20.00%
Mean IoU, per-image macro-average	29.53%
Mean IoU, dataset-level accumulated confusion	19.50%

the two region-scale causes, prior-reliance and distribution-shift, are a small minority of total training pixels with no class-balancing correction. An unweighted
445 classifier trained under this imbalance is expected to converge toward something close to the majority-class rule, which is what Table 2 shows. Figure 2 makes this concrete with one held-out example per procedure: the predicted cause map for the prior-reliance example does recover the general location of the true suppressed region, but with substantial noise and incomplete coverage, exactly
450 the signature of a model that has learned a real but weak signal on top of a strong class prior rather than either pure memorization of the prior or clean per-pixel discrimination. Separately, 40 training images is a small fraction of the design’s target scale, and the training loss curves (kernel estimator: 1.63 dropping to a 0.87–0.93 plateau after one epoch; fusion head: 1.98 dropping
455 to a 1.75–1.76 plateau after roughly eight epochs) both show real but modest learning consistent with a data-limited regime rather than a broken training signal. Section 5 treats class-balanced loss weighting, not just benchmark scale, as an immediate next step rather than optional future polish.

4.4. Disentanglement

460 Table 3 reports the Pearson correlation matrix among the four raw signal maps, flattened and pooled across all pixels of all 20 held-out images, to report

the correlation directly rather than assume clean separation among the four causes.

Table 3: Pearson correlation matrix among S1–S4 on the held-out evaluation set.

	S1 (null-space)	S2 (prior-reliance)	S3 (mismatch)	S4 (distribution shift)
S1	1.000	0.267	−0.060	−0.047
S2	0.267	1.000	−0.106	−0.123
S3	−0.060	−0.106	1.000	0.152
S4	−0.047	−0.123	0.152	1.000

Three of the six off-diagonal correlations are weak, at or below $|r| = 0.15$,
465 indicating the corresponding signal pairs behave close to independently across
the evaluation set. The exception is S1 and S2, correlated at $r = 0.267$: regions
with high null-space uncertainty also tend to show elevated prior reliance. This
is not a defect in the two signals; it is the expected consequence of what each
signal actually measures. A region with very little constraining evidence (high
470 S1) leaves the model with correspondingly less real signal to weigh against its
prior in the first place, which mechanically pushes prior reliance upward through
the same channel that produces the uncertainty. This coupling was flagged as
a risk before any data was collected; this experiment is the first evidence it is
real, and roughly the size the intuition predicted, rather than negligible or so
475 severe as to make the two signals redundant.

4.5. Ablation

Table 4 reports the ablation the evaluation protocol specifies: the trained
fusion head against each raw signal used alone and against an equal-weight com-
bination rule. Concretely, each single-signal baseline predicts that signal’s own
480 cause wherever it exceeds a threshold and reliable otherwise; the equal-weight
rule is the same idea applied jointly, predicting whichever of S1–S4 is largest at
a pixel if that maximum exceeds a shared threshold, reliable otherwise, since
it has no learned weights to combine the four signals any other way. All five

thresholds were grid-searched over 51 evenly spaced points from 0 to 1 on the
485 fusion head’s own training set, then applied frozen to the held-out set, so the
comparison is never fit and evaluated on the same data. This calibration ob-
jective (raw accuracy under severe class imbalance) is itself a factor in what
follows: a threshold that fires only on a rare cause typically costs more false
positives against the dominant reliable class than it gains in true positives, so
490 an accuracy-maximizing grid search is biased toward never firing regardless of
how informative the underlying signal actually is, and the results below should
be read with that bias in mind rather than as a clean measurement of signal in-
formativeness alone. The fusion head’s own accuracy here (48.36%) reproduces
Table 2’s number exactly; its mIoU (29.81%) differs slightly from Table 2’s
495 29.53% despite both coming from the same checkpoints and the same held-out
set, a small discrepancy attributed to ordinary run-to-run floating-point non-
determinism in fp16 diffusion sampling across two separate script executions
rather than to any difference in method, and reported here rather than silently
rounded away.

Table 4: Ablation: fusion head against every raw-signal and hand-combination baseline, on
the held-out set (per-image macro-average).

Method	Pixel accuracy	Mean IoU
Fusion head (ours)	48.36%	29.81%
Equal-weight combination (threshold 1.00)	47.66%	23.83%
S1 (ill-posedness) alone (threshold 1.00)	47.66%	23.83%
S2 (prior-reliance) alone (threshold 1.00)	47.66%	23.83%
S3 (mismatch) alone (threshold 0.24)	46.52%	23.89%
S4 (distribution-shift) alone (threshold 0.56)	47.66%	23.83%

500 Three of the five baselines, equal-weight, S1 alone, and S2 alone, calibrate
to a threshold of exactly 1.0, the upper boundary of the 51-point search grid
rather than an interior optimum, meaning the grid search found that never firing
(predicting reliable at every pixel) maximizes training-set accuracy at every grid

point it tried. S4 alone calibrates to a genuine interior value (0.56), but S4’s own
 505 signal values on the held-out set never reach even that threshold, so it too never
 fires in practice; its accuracy and mIoU end up identical to the three threshold-
 1.0 rules for that reason, not because its threshold is itself degenerate. All four
 therefore match the majority-class baseline in Table 2 exactly, because on this
 data they are the same rule. Only S3 finds a threshold (0.24) that actually
 510 fires, and it does slightly worse than the majority rule on accuracy (46.52%
 vs. 47.66%, a 1.14-point absolute or 2.4% relative drop) while matching it on
 mIoU. The fusion head is the only method in this comparison that exceeds the
 majority-class rule on both metrics: by 0.70 points (1.5% relative) on accuracy,
 and by 5.98 points (25.1% relative) on mIoU, reporting both conventions here so
 515 neither metric’s margin is presented more favorably than the other. The larger
 mIoU margin has an identifiable mechanical cause, not a different underlying
 effect: mIoU credits correctly identifying minority-class pixels (prior-reliance,
 distribution-shift) that every degenerate baseline scores zero intersection on by
 never predicting them at all, so accuracy, which is dominated by the numerically
 520 large reliable and whole-image classes, is far less sensitive to this difference than
 mIoU is.

This comparison has two limits. First, it rules out five specific hand-crafted
 alternatives, not hand-crafted alternatives in general: a jointly-fit but still simple
 classifier on the same four raw signals, such as a per-pixel logistic regression,
 525 sits in the unexplored middle ground between a single grid-searched threshold
 and the full convolutional fusion head, and this pilot does not test it. Second,
 the margins above come from one run on 20 held-out images with no per-image
 variance or significance test reported, so they should be read as this pilot’s point
 estimate rather than as a result shown to be robust to resampling. Both are
 530 concrete, low-cost additions for the next iteration of this evaluation rather than
 open-ended future work. With those two caveats attached, this paper reads the
 result as evidence the fusion head has learned something beyond the class prior
 among the alternatives actually tested here, not as evidence that no simpler
 method could ever match it: the fusion head has not been shown to be a strong

535 classifier at this pilot’s scale, but it has been shown to be the only method among those tested capable of attempting minority-class attribution at all, which is a real, if modest, form of value the raw signals and their naive combination do not provide on their own.




Figure 2: One held-out example per xSR-CausalBench procedure. LR input (nearest-neighbor upsampled for display), the backbone’s $16\times$ SR output, the ground-truth cause map, and the fusion head’s predicted cause map. Color key: ill-posed (red), prior-reliance (green), degradation-mismatch (yellow), distribution-shift (blue), reliable (gray). The prior-reliance row shows the predicted map recovering the approximate location of the true suppressed region with visible noise, consistent with the weak-but-real signal described above; the whole-image ill-posed and mismatch rows are largely uniform in both ground truth and prediction by construction.

5. Discussion and Limitations

5.1. *What this pilot validates, and what it does not*

The experiment reported here demonstrates that the CHASR architecture trains end to end on real images and produces a fusion head whose predictions are comfortably above uniform-random and, in the one qualitative example examined directly (Figure 2), spatially correlated with the true region rather than uniform noise; it also demonstrates that the disentanglement analysis specified in advance produces an interpretable, honest result rather than an embarrassing surprise. It does not yet demonstrate that the fusion head has learned substantially more than the benchmark’s own class prior: the 0.7-point margin over the majority-class baseline (Table 2) is the single most important number in this paper to improve before any stronger claim is warranted, and framing the pixel-accuracy result as state-of-the-art or as clearly beyond the class prior would misrepresent both the scale of the experiment and what the numbers can currently support. The training and evaluation sets used here, 40 and 20 images respectively, are roughly one twentieth of the design’s target scale of 500 to 1000 images per procedure. Scaling to that target, incorporating Urban100 and Manga109 alongside DIV2K for structural and textural diversity as the design specifies, is the immediate and clearly defined next step, not a vague direction for future work.

One further evaluation component the design specifies was not run in this pilot and is named explicitly rather than left implicit: the downstream case study correlating flagged high-mismatch or high-distribution-shift regions with actual optical character recognition failure on forensic-style text crops, which ties most directly to this special issue’s task-driven framing, was implemented as a reusable module but not exercised against real recognition failures in this pilot. This changes what can be claimed about CHASR’s practical value and does not require a change to the architecture itself; it is evaluation work, not redesign work.

Section 4.5 supplies the ablation this same protocol calls for, comparing

the fusion head against each signal alone and against an equal-weight hand-
 570 combination rule. Its result sharpens, rather than resolves, the 0.70-point
 margin over the majority-class baseline flagged above (Table 2): none of the
 five hand-crafted alternatives tested outperforms the trivial majority-class rule,
 while the fusion head is the only one that does, with a substantially larger mar-
 575 head has learned something beyond the class prior among the specific alterna-
 tives tested, not as evidence that no simpler or differently-calibrated method
 ever could; Section 4.5 names the two most direct ways that stronger claim
 remains open (a jointly-fit simple classifier on the same signals, and per-image
 variance across a larger held-out set) rather than treating the ablation as closing
 580 the question.

Two measurement choices also warrant flagging directly. Signal 1’s variance
 estimate is computed from $K = 2$ stochastic samples per pixel, a very small
 sample for a variance estimate; the qualitative pattern it produces is consistent
 with expectation, but the estimate itself is noisy at this K , and raising it is cheap
 585 (it costs additional backbone forward passes, not additional training) compared
 to every other scaling step this section names. The mean IoU reported in Table 2
 is a per-image macro-average, the mean of one mean_IoU computation per held-
 out image, rather than the dataset-level convention common in segmentation
 literature, which accumulates a single confusion matrix over the entire held-out
 590 set before dividing per class; the two are not numerically equivalent, and Table 2
 reports the dataset-level number alongside it so this pilot’s mIoU is not read as
 more standardized than it is.

5.2. Factor entanglement

The $r = 0.267$ correlation between null-space uncertainty and prior reliance,
 595 reported without qualification in Section 4, means the four-way attribution this
 paper proposes should not be read as claiming four cleanly independent causes
 that a classifier simply sorts a pixel into. It is closer to four correlated but distin-
 guishable explanatory factors, in the same sense that a set of correlated predic-

tors in a regression can each still carry real, separable information even though
 600 they are not orthogonal. The practical consequence for a downstream user is
 that a pixel attributed primarily to prior reliance in a region that also shows
 high null-space uncertainty should be read as “the model is filling this in largely
 from its prior, and there was little evidence available for it to do otherwise,” a
 compound but coherent statement, rather than as two unrelated observations
 605 that happen to coincide. Whether this entanglement grows, shrinks, or restruc-
 tures at the design’s full target scale is an open, empirical question this pilot’s
 sample size cannot answer with confidence, and is worth re-measuring once the
 benchmark is scaled rather than assumed to hold at the same magnitude.

5.3. *Synthetic-to-real gap*

610 Every ground-truth label used to train and evaluate the fusion head comes
 from xSR-CausalBench’s controlled synthetic construction, because the true
 cause of a hallucination in a genuine photograph is, by definition, not ob-
 servable. HalluGen’s controllable-hallucination construction and the Halluci-
 nation Score’s severity judgments face the same constraint; no hallucination-
 615 attribution method plausibly avoids it. What can and should be strengthened
 is validating that performance on the synthetic benchmark transfers to genuinely
 uncontrolled degradations, using real low-resolution/high-resolution pairs from
 datasets such as RealSR or DRealSR where no synthetic cause label exists but
 severity correlation and qualitative inspection remain possible. This paper’s
 620 pilot evaluates on synthetic pairs only; the real-world generalization question
 remains open.

5.4. *Backbone dependency and forensic-use caveats*

CHASR’s four signals are computed against one specific frozen backbone,
 and both the calibration of Signal 3’s mismatch threshold and Signal 4’s ref-
 625 erence feature bank are tied to that backbone’s specific training distribution.
 Swapping backbones, or fine-tuning the existing one, would require rebuild-
 ing the feature bank and likely re-tuning the mismatch calibration, though the

four-signal architecture and the fusion head training procedure would carry over unchanged. Readers considering CHASR or a similar attribution system for genuinely forensic use should also note that a calibrated reliability score reduces but does not eliminate the risk of over-trusting a hallucinated region: the fusion head’s reliability output is itself a learned prediction, trained on synthetic labels at pilot scale, and should be treated as an additional, imperfect signal to weigh alongside domain expertise rather than as a certification that a flagged-reliable pixel is ground truth.

6. Conclusion

Extreme super-resolution’s central practical problem is not that it hallucinates, since some degree of hallucination is an unavoidable consequence of the information the degradation destroys, but that existing systems give a user no way to distinguish an inherently unrecoverable region from a model that ignored evidence it should have used, from a mismatched operating condition, from genuinely unfamiliar content. This paper introduced a four-way causal taxonomy for exactly this distinction, the CHASR architecture that computes it as four signals fused by a small trained head sitting on top of an otherwise frozen backbone, and xSR-CausalBench, a synthetic benchmark that makes the four causes separately trainable and evaluable. A first end-to-end validation at real $16\times$ upscaling shows the architecture trains successfully and produces attribution accuracy well above uniform-random, though only narrowly above the benchmark’s own class prior at this pilot’s training scale, together with a disentanglement analysis that confirms, rather than merely speculates about, a factor-entanglement risk flagged before any data was collected, and an ablation that shows the trained fusion head, unlike every hand-crafted signal-threshold or equal-weight alternative tested, exceeds the benchmark’s own majority-class rule on both accuracy and mean IoU. The path from this pilot to a complete evaluation, class-balanced training at larger benchmark scale, and the downstream task-driven case study this special issue is centrally concerned with, is

concrete and already scoped, and is where this line of work goes next.

Data Availability Statement

660 The implementation, including the degradation pipelines, the four signal estimators, the fusion head, xSR-CausalBench construction code, and the training and evaluation scripts used to produce every number reported in this paper, is available at https://github.com/vinhqdang/super_resolution. DIV2K [20] is a publicly available dataset used under its original terms.

Ethics Declaration

665 This study used only publicly available image data (DIV2K) and did not involve human subjects, personal data, or any content requiring ethics board review.

Author Contributions (CRediT)

670 Q.D.: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review & editing.

Declaration of Competing Interest

The author declares no competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

675 Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used Claude (Anthropic) and Claude Code in order to assist with software implementation, experiment execution, and drafting and revising the manuscript text. After using this tool, the author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

References

- [1] S. Kim, H. F. J. Tregidgo, C. Jin, M. Figini, D. C. Alexander, Hallugen: Synthesizing realistic and controllable hallucinations for evaluating image restoration, arXiv preprint arXiv:2512.03345 (2025).
- [2] W. Ren, R. Goyal, Z. Hu, T. T. Aumentado-Armstrong, I. Mohomed, A. Levinshtein, Hallucination score: Towards mitigating hallucinations in generative image super-resolution, arXiv preprint arXiv:2507.14367 (2025).
- [3] S. Bell-Kligler, A. Shocher, M. Irani, Blind super-resolution kernel estimation using an internal-gan, in: Advances in Neural Information Processing Systems (NeurIPS), 2019.
- [4] J. Liang, G. Sun, K. Zhang, L. Van Gool, R. Timofte, Mutual affine network for spatially variant kernel estimation in blind image super-resolution, in: IEEE/CVF International Conference on Computer Vision (ICCV), 2021.
- [5] L. Wang, Y. Wang, X. Dong, Q. Xu, J. Yang, W. An, Y. Guo, Un-supervised degradation representation learning for blind super-resolution, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2021.
- [6] J. Wang, Z. Yue, S. Zhou, K. C. Chan, C. C. Loy, Exploiting diffusion prior for real-world image super-resolution, International Journal of Computer Vision (2024).

- [7] R. Wu, T. Yang, L. Sun, Z. Zhang, S. Li, L. Zhang, SeeSR: Towards semantics-aware real-world image super-resolution, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024.
- [8] T. Yang, R. Wu, P. Ren, X. Xie, L. Zhang, Pixel-aware stable diffusion for realistic image super-resolution and personalized stylization, in: European Conference on Computer Vision (ECCV), 2024.
- [9] B. S. Kim, J. Kim, J. C. Ye, Chain-of-zoom: Extreme super-resolution via scale autoregression and preference alignment, in: Advances in Neural Information Processing Systems (NeurIPS), 2025, spotlight presentation. arXiv:2505.18600.
- [10] X. Wang, L. Xie, C. Dong, Y. Shan, Real-ESRGAN: Training real-world blind super-resolution with pure synthetic data, in: IEEE/CVF International Conference on Computer Vision Workshops (ICCVW), 2021.
- [11] R. Wang, Y.-Y. Chuang, Uncertainty-guided latent diffusion models for faithful super resolution, arXiv preprint arXiv:2608.25998 (2026).
- [12] J. Yin, J. Li, H. Xing, QUSR: Quality-aware and uncertainty-guided image super-resolution diffusion model, in: IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2026, arXiv:2603.09125.
- [13] K. N. H. Vo, A. Vazhentsev, A. Shelmanov, T. Baldwin, Y. Kementchedzhieva, Failing to see or failing to know? attributing errors in vision-language models, arXiv preprint arXiv:2607.04683 (2026).
- [14] Q.-V. Dang, N.-S.-A. Nguyen, T.-B.-D. Vo, M. N. Dinh, A. D. Le, H. V. Vu, P. L. Nguyen, Real-time hallucination correction in vision-language models using dynamic knowledge graph verification, Discover Artificial Intelligence (2026). doi:10.1007/s44163-026-01778-z.

- [15] S. K. Aithal, P. Maini, Z. C. Lipton, J. Z. Kolter, Understanding hallucinations in diffusion models through mode interpolation, in: Advances in Neural Information Processing Systems (NeurIPS), 2024.
- 735 [16] Y. Wang, J. Yu, J. Zhang, Zero-shot image restoration using denoising diffusion null-space model, in: International Conference on Learning Representations (ICLR), 2023, oral presentation, notable-top-25%.
- [17] A. Sendjasni, M.-C. Larabi, Embedding similarity guided license plate super resolution, arXiv preprint arXiv:2501.01483 (2025).
- 740 [18] K. S. Krishnan, K. S. Krishnan, HAT super-resolution and a PARSeq+CLIP4STR voting ensemble for extreme in-the-wild license plate recognition, arXiv preprint arXiv:2607.08896 (finalist entry, XLPSR Grand Challenge, ICIIP 2026) (2026).
- 745 [19] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, et al., DINOv2: Learning robust visual features without supervision, Transactions on Machine Learning Research (2024).
- [20] E. Agustsson, R. Timofte, NTIRE 2017 challenge on single image super-resolution: Dataset and study, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 2017.
- 750 [21] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, O. Wang, The unreasonable effectiveness of deep features as a perceptual metric, in: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2018.